

Effizientes Programmieren in C, C++ und Rust

Intro and C Basics

Colin Bretl, Nikolai Maas, Peter Maucher | 24.10.2024




Table of Contents

1. Program Efficiently?
2. Our Performance Model
3. (Instruction) Caching
4. (Instruction) Prefetching
5. C Introduction
6. Recap

Program Efficiently?

Why program efficiently?

- Be *FAST*
- But why is fast important?

Why fast?

Discuss!

ENABLE SOLUTIONS

- Solve large datasets (e.g., CERN)
- Overcome hardware limits (your laptop)

SAVE MONEY

- Less hardware
- Less CPU time (AWS, ...)
- Less power (environmental impact)
- Time of employees

IMPROVE (CUSTOMER) EXPERIENCE

- Shorter response time
- Less wait time
- All of the above

How to be fast?

OTHER LECTURES

- Algorithms \Rightarrow (mostly) other lectures
- Data structures \Rightarrow (somewhat) other lectures
- Design for Performance \Rightarrow (somewhat) other lectures
- Parallelization \Rightarrow (mostly) other lectures

THIS LECTURE

- Data layout
- Optimizations
 - Machine characteristics
 - Caches
 - Prefetchers
 - ...
 - Code structure
 - Loop unrolling
 - Dynamic Dispatch vs Generic Code
 - ...
- Profiling
- Data modelling

Prior Knowledge

We assume knowledge in:

- *Operating system* ↑
- *Basic algorithms* ↑
- *Computer architecture* ↑

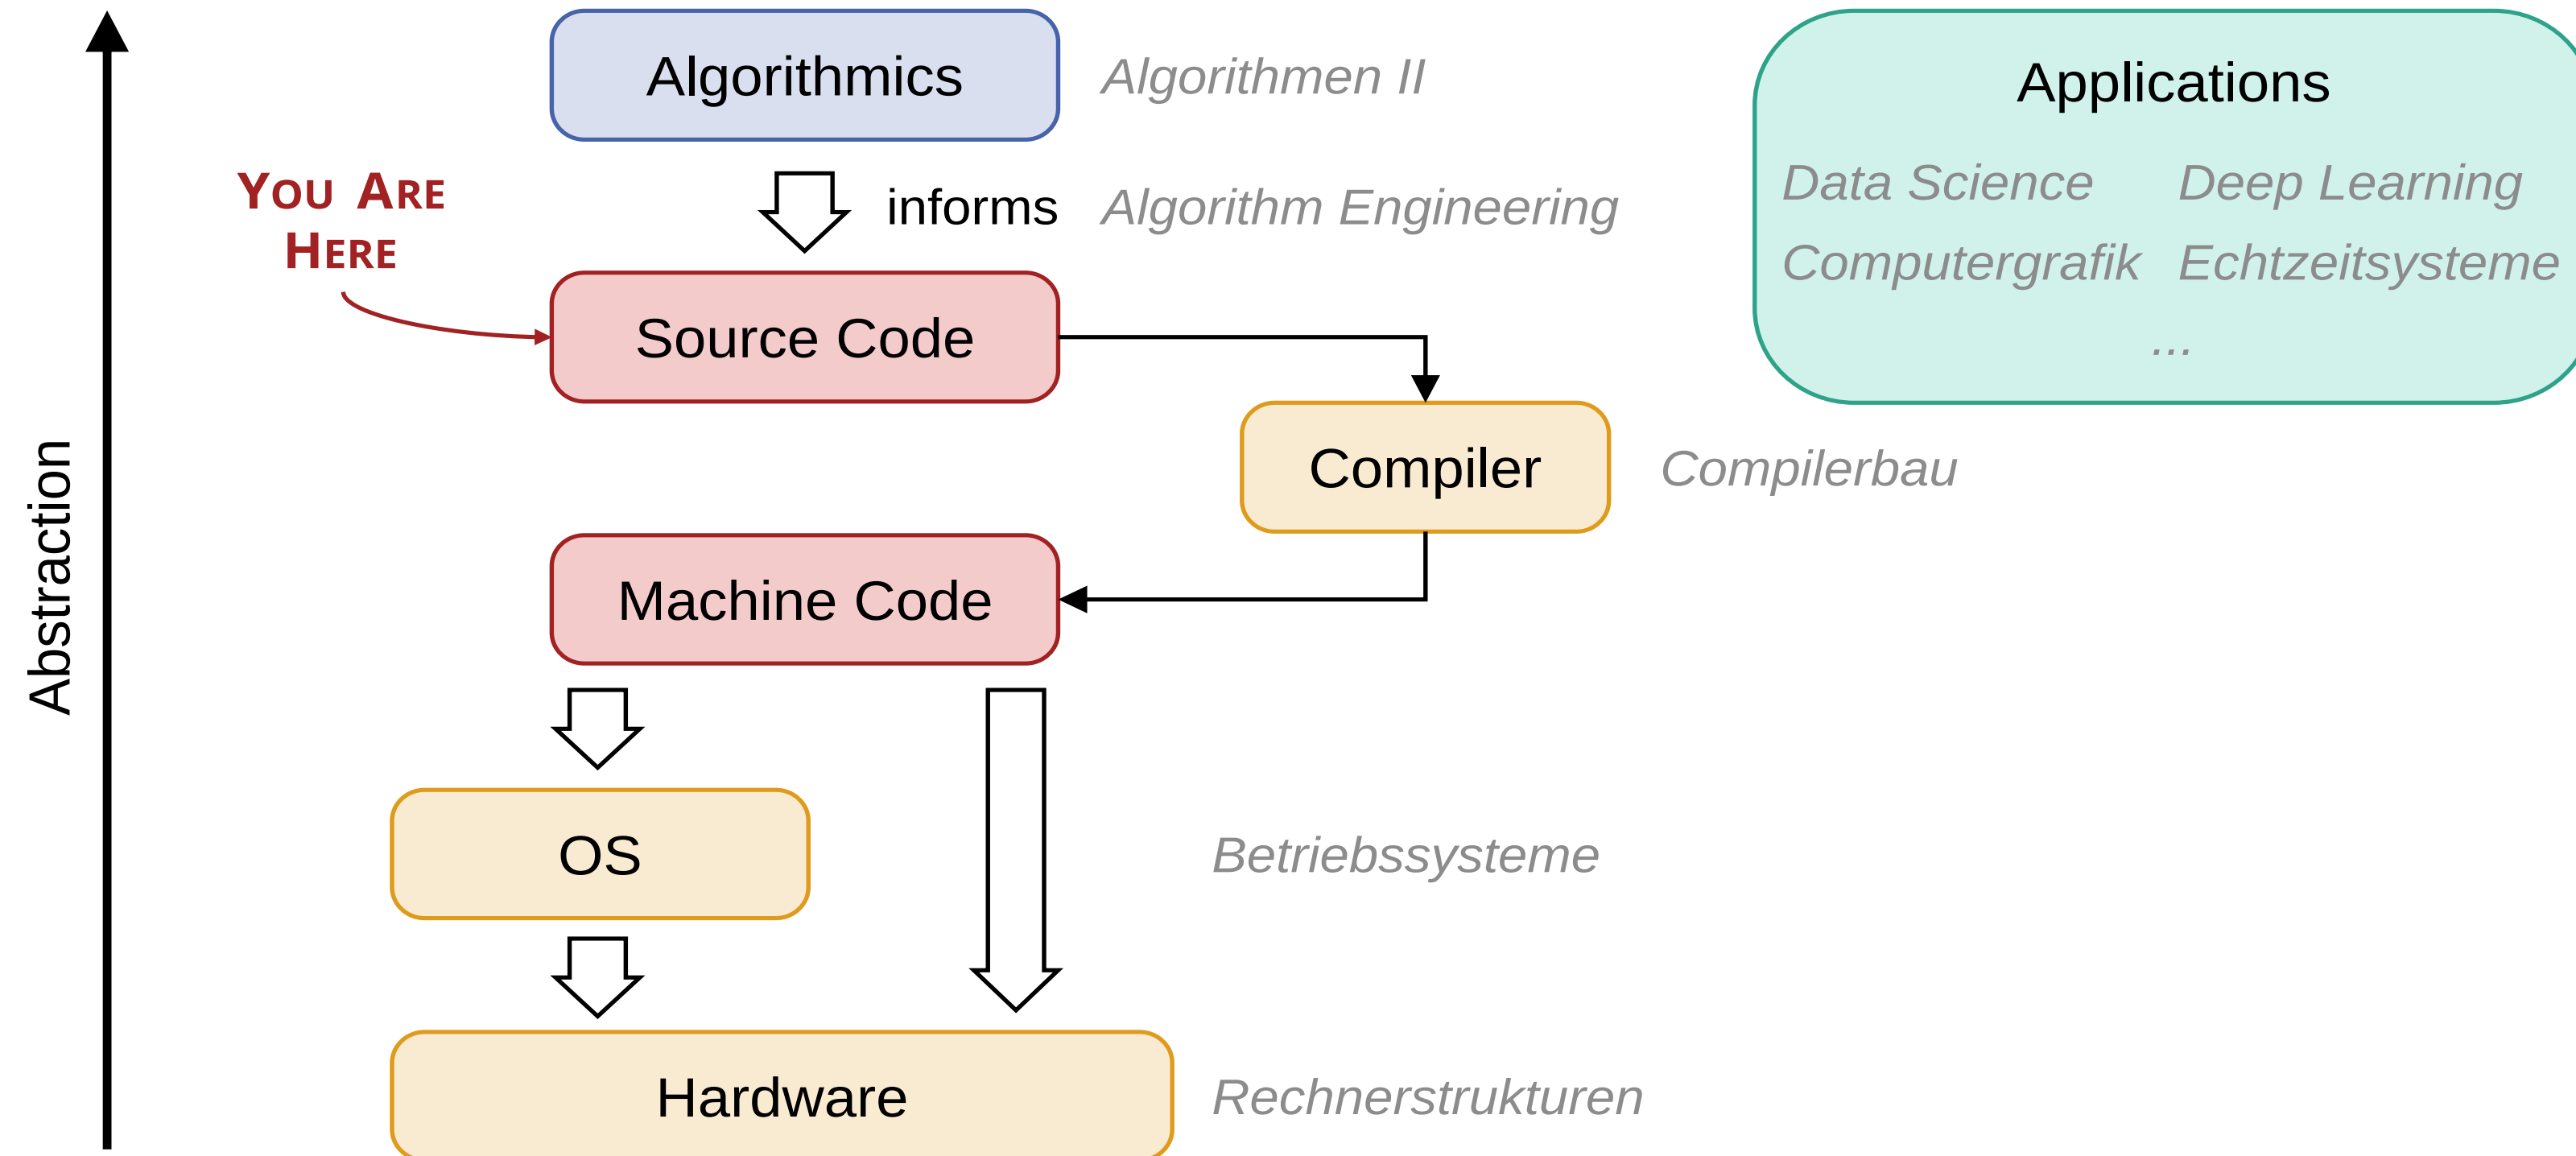
We do not assume knowledge in:

- Advanced computer architecture ↗
 - See next slides...
- Advanced algorithms ↗
 - Introduced as needed

Assumed Knowledge

- *Operating system* ↑
 - We (mostly) do not inspect OS
 - Programs start at `main()` and end at `exit()`
 - We ignore implementation details, like drivers
 - System libraries exist
 - Historical baggage is just that
 - we *DO* model OS performance ↗
- *Basic algorithms* ↑
 - Big O notation
 - Some basic algorithms as presented in GBI, Algo 1
- *Computer architecture* ↑
 - No lecture referencable
 - Need more than "Rechnerorganisation"
 - Assumed: 5-Stage RISC Pipeline
 - Caches
 - Some assembly knowledge (x86, RISC-V, Pseudoops)

The Context of this Lecture



There will be frequent backward- ↖, forward- ↗, and external- ↑ references

Our Performance Model

Or How I Learned to Stop Worrying and Love the Prefetcher

RO: 5 STAGE DLX ("DELUXE") PIPELINE

the DLX [DeLuXe] 5-Stage RISC pipeline

Reminder: The 5-Stage RISC pipeline

- In order \Rightarrow one slow instruction slows down all following instructions
- Stall on memory access
- (Relatively) low clock
- Weird hacks (branch delay slot \Rightarrow nops everywhere)
- Uncore (TLB, Caches, ...) not discussed

Our Performance Model

Or How I Learned to Stop Worrying and Love the Prefetcher

ACTUAL CONTEMPORARY

HW: AMD ZEN 4

- Out of Order (OoO)
- A lot of functional blocks
- Extremely fast
- Multiple layers of optimizations
- Difficult to understand
- Sometimes weird performance characteristics
- Not discussed in detail in this lecture

Zen 4 block diagram

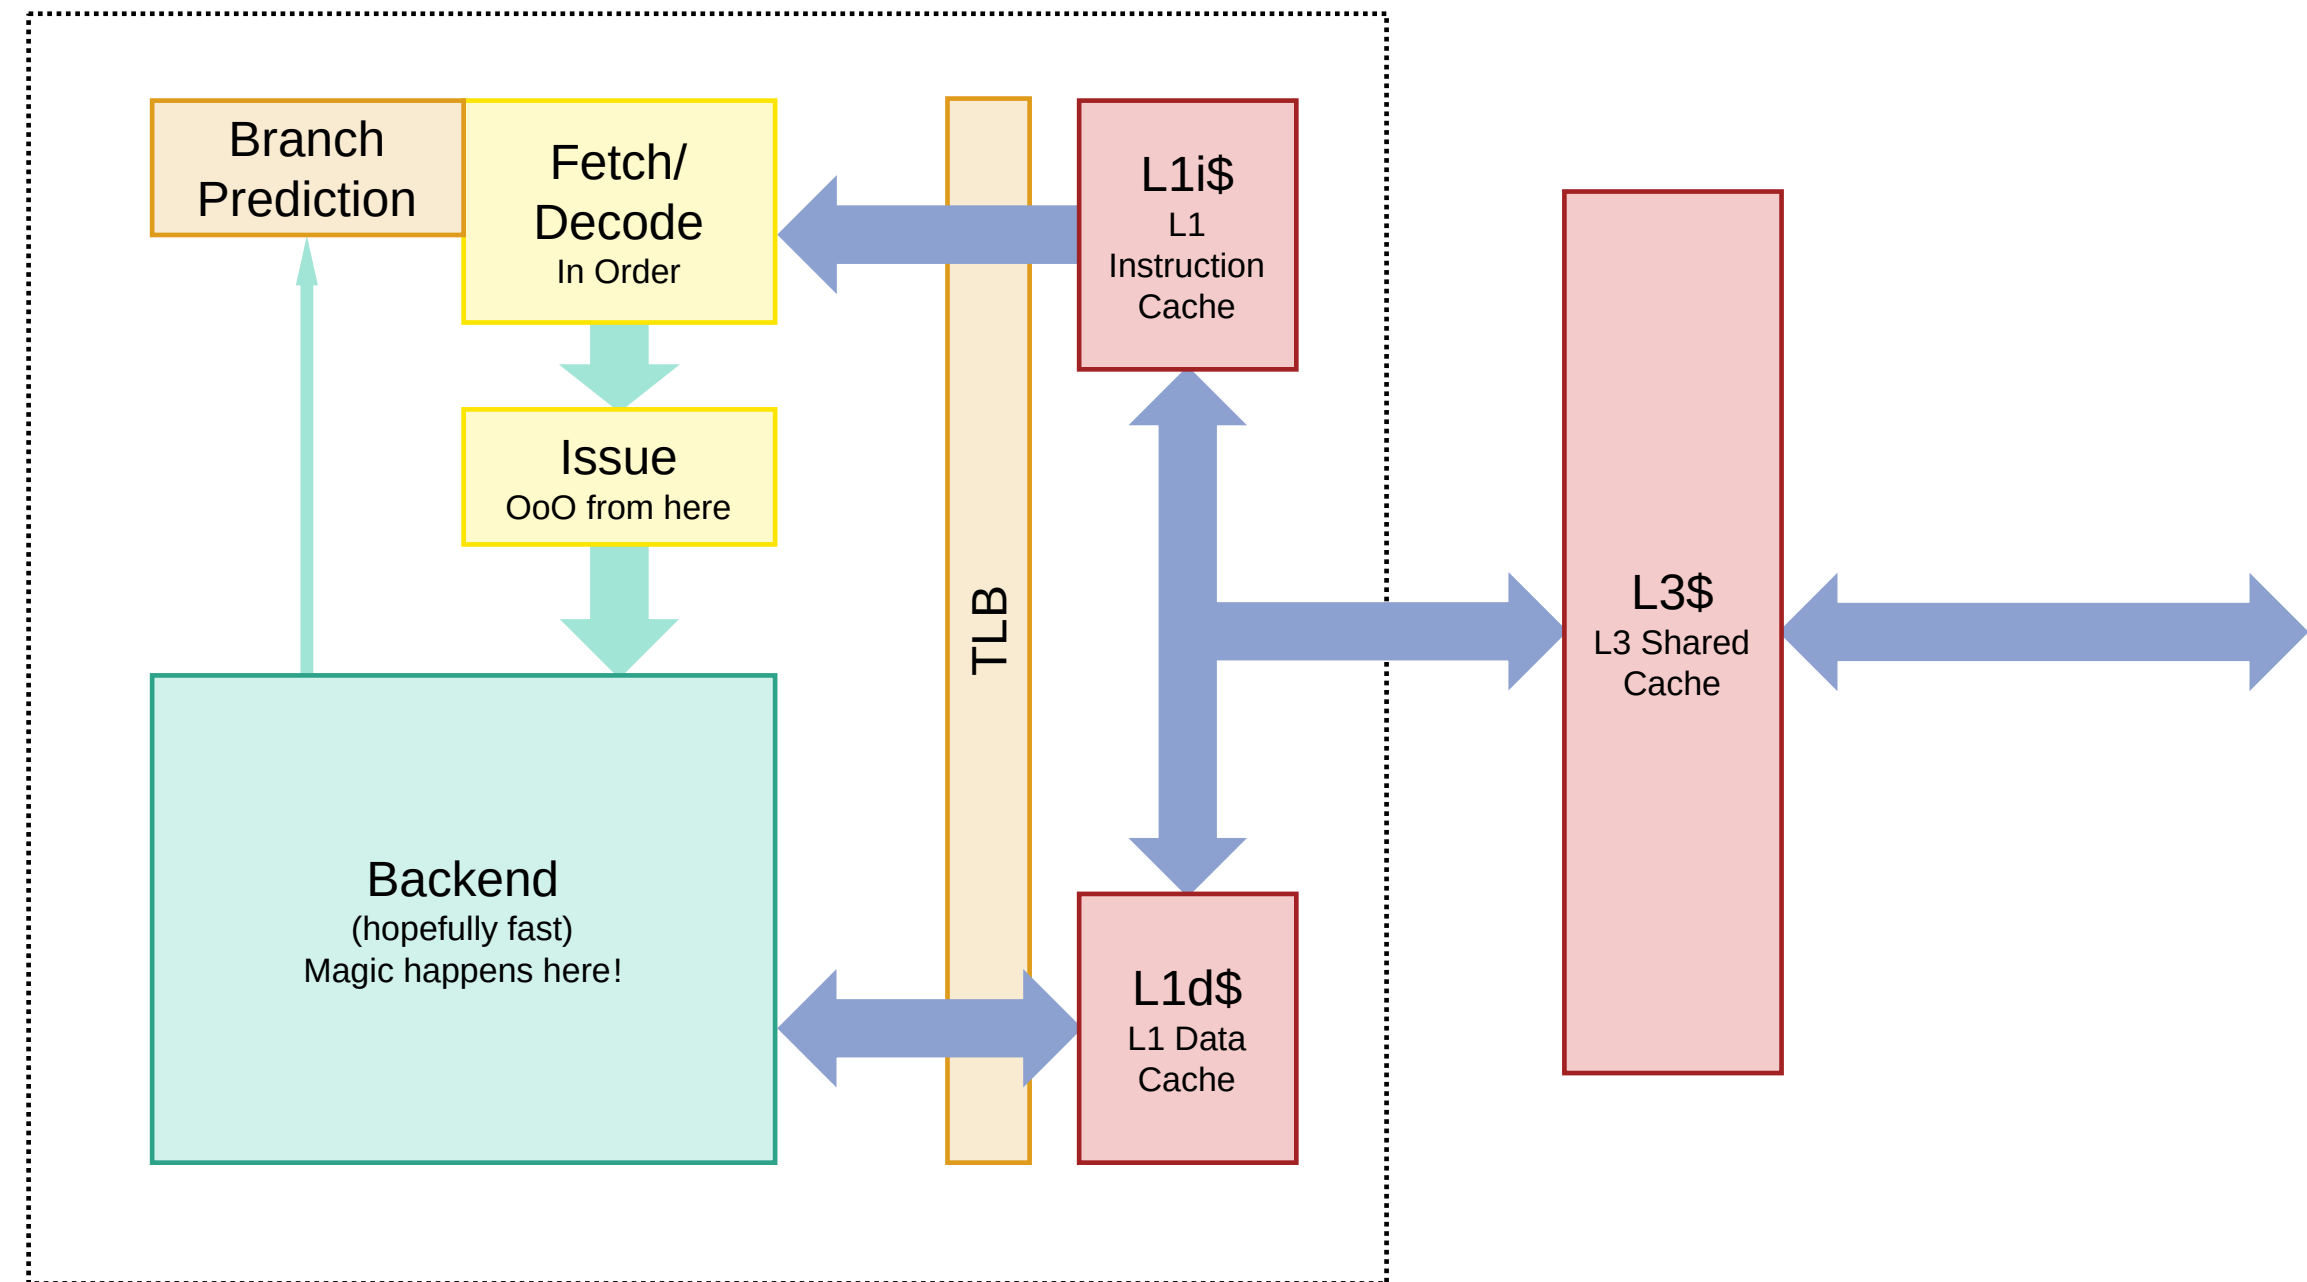
Real world: AMD Zen 4

Our Performance Model

Or How I Learned to Stop Worrying and Love the Prefetcher

THIS LECTURE: IN-BETWEEN & BLACK BOXED \Rightarrow INTUITION

- Simple OoO intuition, OoO on instruction level
- simple Branch predictor
- Only L1 and L3 cache
- Single-level TLB
- Linear prefetcher
- No forwarding \Rightarrow memory latency either L1 or L3 or DRAM



This lecture: μ EffCpp

(Instruction) Caching

Money, Money, Money

```
int sum = 0;
for(int i = 0; i < 3; ++i) sum += i;
```

```
addi a0, zero, 0 # sum
addi t0, zero, 0 # i
addi t1, zero, 3 # limit
loop:
    beq t0, t1, end
    add a0, a0, t0
    addi t0, t0, 1
    jal zero, loop
end:
```

How long does it take? (rough estimation)

- Completely executed ("retired") Instructions: 3 Preamble + 3 × 4 Loop Body + 1 Postamble = 16 Instructions
- Instruction execution: about 1 cycle (4 GHz: .25 ns) ⇒ 4 ns
- But wait, there's more!

- Issue: instruction fetch: about 75 ns (300 × slower!) (\sum 1204 ns)
- Idea: cache loaded instruction ⇒ 7 Fetches, 16 Instructions = 529 ns (2.28 × speedup!)
- Best case: Fetch all at once, 75 + 4 = 79 ns, 15,25x speedup

(Instruction) Caching

Cache Lines

```
int sum = 0;
for(int i = 0; i < 3; ++i) sum += i;
```

```
addi a0, zero, 0 # sum
addi t0, zero, 0 # i
addi t1, zero, 3 # limit
loop:
    beq t0, t1, end
    add a0, a0, t0
    addi t0, t0, 1
    jal zero, loop
end:
```

Best case last slide:

Fetch all at once, $75 + 4 = 79$ ns, 15,25x speedup

- Number of bytes to fetch?
RISC-V: 4 Byte per Instruction: 7×4 Bytes = 28 Bytes
- Number of bytes per memory transaction?
(usually) 64 bytes
- Number of fetches required? sometimes 2! (often 1)

What does number of fetches depend on?

- Memory organized as cache lines
- Cache line (usually) size of memory transaction
- Cache lines are aligned:
 $[n \times 64, n \times 64 + 63], n \in \mathbb{N}$
- Crossing boundaries: Both lines need to be fetched

(Instruction) Caching

Speedup to the moon! 🏍️

```
int sum = 0;
for(int i = 0; i < 1000; ++i) sum += i;
```

```
addi a0, zero, 0 # sum
addi t0, zero, 0 # i
addi t1, zero, 1000 # limit
loop:
    beq t0, t1, end
    add a0, a0, t0
    addi t0, t0, 1
    jal zero, loop
end:
```

- Retired Instructions: 3 Preamble + 1000×4 Loop Body + 1 Postamble = 4004 Instructions
- Cacheless execution: $\Rightarrow 4004 * (0.25 + 75) = 1001 \text{ ns} + 300300 \text{ ns} = 301301 \text{ ns}$
- Cached Execution: $1001 \text{ ns} + 75 \text{ ns} = 1076 \text{ ns}$
- Speedup: $280 \times$
- Limit? $75.25 \text{ ns} / .25 \text{ ns} = 301 \times$

(Instruction) Caching

A (Real) Cache Hierarchy

Zen 4

Level	Name	Latency	Factor
"0"	Registers	1 cycle (often even lower)	<i>n/a</i>
L1	L1 Cache	4 cycles (< 1 ns)	$\geq 4 \times$
L2	L2 Cache	14 cycles (≈ 2.5 ns)	$3.5 \times$
L3	L3 Cache	≈ 8 -9 ns	$\approx 3.5 \times$
"4"	(D) RAM	≈ 75 ns	$\approx 8.8 \times$

- Model:
Registers $\times 4 \Rightarrow$
L1 $\times 12.25 \Rightarrow$
L3 $\times 8.9 \Rightarrow$ RAM
- L1 Latency: 4 cycles...
- Assumption before: Instruction fetch in "0" cycles?
- Solution: Prefetching

(Instruction) Prefetching

Gimme, Gimme, Gimme

[illegible]

- Is that efficient assembly? No
- But does it show my point? Yes
- What is the next instruction? The one below
- **Idea prefetching:** Preload data that is expected to be used into cache / core
- **Our Model:** Simple prefetcher that only guesses linear patterns
- **Linear pattern:** Prefetch the data at same offset of previous two loads, if "enough" such loads have happened

DISCLAIMERS AND PRELIMINARIES

This lecture is opinionated :)

Why **C**?

Language to implement **Unix** in

- 👍: No assembly!
- 👍: Generic memory management!
- 👍: Platform-independent code!
- 👍: Simple compiler!
- 👍: Fast!
- 👍: Library writeable in **C**!

TL;DR: It was a great tool in 1972...

Why **C** today?

- Mostly for supporting existing projects
- Lingua franca: Interface language for a lot of languages
- Quite influential
- Performance reference
- Ship precompiled? ⇒ **C** an option
- Memory management still important topic

A Simple C Program

Why?

Get out your laptops, it's practice time!

Hello World *(it had to be done)*

```
#include <stdio.h>
void main(int argc, char **argv) {
    printf("%s %s\n", "Hello", "EffCpp");
}
```

And now?

Compile and execute

```
vesemir@kaermorhen ~/effcpp/C1$ gcc hello.c -o hello
vesemir@kaermorhen ~/effcpp/C1$ ./hello
Hello EffCpp
vesemir@kaermorhen ~/effcpp/C1$
```

Wait...

```
vesemir@kaermorhen ~/effcpp/C1$ gcc hello.c -O3 -o hello
vesemir@kaermorhen ~/effcpp/C1$ ./hello
Hello EffCpp, but faster!
vesemir@kaermorhen ~/effcpp/C1$
```

En Détail

```
#include <stdio.h>
void main(int argc, char **argv) {
    printf("%s %s\n", "Hello", "EffCpp");
}
```

```
#include <stdio.h>
```

- Used to include other files ↗
- Here: what is this `printf()`?

```
void main(int argc, char **argv) {
```

- The program begins here
- `argc`, `argv` command line parameters

```
    printf("%s %s\n", "Hello", "EffCpp");
```

- print to stdout
- Format string: see manpage ↑

```
}
```

- This is the end *my only friend, the end*
- Closes function ↗, block ↗, and scope ↗

A Simple C Program

Say Hello to Java

```
int friends() {
```

It was raining **while** you waited **for** the plane to arrive. But what **else** could you **do**? In any **case**, it was no question **if** they arrived, the pilots aren't taking a **break**, they'd **continue** the approach. You already planned where y'all would **goto** ↗ after the airport, and then **switch** to another club. You're very happy they all did **return** EXIT_SUCCESS;

```
}
```


Control Flow

Oh Look, a Performance Discussion!

WHICH ALTERNATIVE IS BETTER?

```
int sum = 0;
for(int i = 0; i < 3; ++i) sum += i;
```

```
int sum = 0;
sum += 1;
sum += 2;
```

AND NOW?

```
int sum = 0;
for(int i = 0; i < 4096; ++i) sum += i;
```

```
int sum = 0;
sum += 1;
...
sum += 4095;
```

- Few instructions \Rightarrow good for i-cache
- Bad: branches
- Good to predict
- Easy to understand

- Good: No branches: good for GPU (?)
- Bad: cache size
- Bad: hard to understand, easy to mess up

And what is better? \Rightarrow Depends, go benchmark! ↖

PS: Why not Gauss? Example also valid for other cases, also compiler will optimize.

NUMBERS

- Unsized signed integers
 - `signed char` ↗
 - `short`
 - `int`
 - `long`
 - `long long`
- Unsized floating point numbers
 - `float`
 - `double`
 - `long double`

NOT NUMBERS

- `void` ↗
- `char` ↗
- `_Bool` *Ab C23:* `bool`
- Functions ↗
- `structures` ↗
- Arrays ↗
- Pointers ↗
- `unions` ↗
- `enumerations` ↗

Types

Numbers

Remember platform independency? Data width: kind of important

- `short`: ≥ 16 bit
- `int` ≥ 16 bit, *commonly* 32 bit
- `long` ≥ 32 bit, *no consensus!*
- `long long` ≥ 64 bit
- `char`: Equivalent to a byte
- `signed char`: Equivalent to a signed byte

enumerations

Numbers, but Named?!

```
if (status == 4) {
    status = 5;
    break;
} else {
    status = 0;
    continue;
}
```

- magic numbers
- often: encode status

```
if (status == SUCCESS) {
    status = COMPLETE;
    break;
} else {
    status = START;
    continue;
}
```

- Clearly readable

USAGE

```
enum enum_name {
    FIRST_CONSTANT,
    SECOND_CONSTANT,
    THIRD_CONSTANT,
    DIRECTLY_ASSIGNED = 7,
    NEGATIVE_NUMBERS = -5,
    SUCCESSOR,
    REASSIGNMENT_IS_VALID = 7
}; // this semicolon is important!
```

- Counting up, starting at zero
- Used without enum name

```
int main() {
    printf("%d %d %d %d %d %d %d\n",
        FIRST_CONSTANT, SECOND_CONSTANT, THIRD_CONSTANT,
        DIRECTLY_ASSIGNED, NEGATIVE_NUMBERS, SUCCESSOR,
        REASSIGNMENT_IS_VALID);
}
```

Result?

0 1 2 7 -5 -4 7

 Performance: Well, it's numbers. Use often!

Functions

They're (mostly) like Java!

- Functions perform calculations that are repeatedly needed
- Functions can have side effects ↗
- Functions have a name (good for code structure)
- Functions take parameters
- Functions return (at most) one value (use `void` for zero)
- There is no Java-`"this"`

```
int add(int first, int second) {  
    for (int i = 0; i < second; ++i)  
        first += 1;  
    return first;  
} // look ma, no semicolon
```

Yes, it's wrong. Here's a fixed version:

```
int add_fixed(int first, int second) {  
    switch(second < 0) {  
        case false:  
            for (int i = 0; i < second; ++i)  
                first += 1;  
        default:  
            return first;  
        case true:  
            for (int i = 0; i > second; --i)  
                first -= 1;  
            return first;  
    }  
}
```

Disclaimer: Code written by trained professionals in a secure environment. Do not try at home.

Write Your Own Type

MOTIVATION

- Often: more information than one number
- Idea: `structure` data, name

Whats not in a C `struct`?

- Code
- Default values
- Inheritance / OOP
- Visibility limits

DEFINITION

```
struct coord {  
    int x;  
    int y;  
}; // this semicolon is (again) important!  
  
struct rect {  
    struct coord upper_left; // structs can contain structs  
    struct coord lower_right;  
};
```

INITIALIZATION

```
int main() {  
    struct coord undefined; // so it begins  
    struct coord zero_initialized = {};  
    struct coord value_initialized = {1, 2};  
    struct coord copy = value_initialized;  
  
    struct rect undef;  
    struct rect zero_init = {};  
    struct rect value_init = {copy, copy};  
    struct rect explicit_def = {{1, 2}, {3, 4}};  
}
```

Bonus Question: Is this valid and what is the result?

```
struct coord coord = {1};
```

Valid! `coord.x == 1` && `coord.y == 0`

Usage

HOW TO USE?

```
struct coord {int x; int y;};
struct rect { struct coord up_left; struct coord low_right;};

struct coord swap(struct coord input) {
    return (struct coord) {input.y, input.x};
}

int main() {
    struct coord undefined;
    undefined.x = 1; // select member with "."
    undefined.y = 2; // now completely defined!
    printf("%d\n", undefined.x); // => "1"

    struct coord init = {4, 5};
    struct rect rect = {init, {7, 8}};
    // "." can be chained for recursive access
    printf("%d %d\n", rect.up_left.x, rect.low_right.y); // => "4 8"
    printf("%d\n", swap(rect.up_left).x); // => "5"
}
```

WHEN TO USE?

- Collect semantically related entities
- Collect values often used together
- Return multiple values from functions
- Name things

WHEN NOT TO USE?

- To pass unrelated parameters to function

What did we learn today?

- What is efficiency and why is that important?
- What do you already know about performance?
- What is our performance model?
- Introduction into C
 - Control Flow
 - Builtin Types
 - `enum`
 - Functions
 - `struct`

And next lecture? ↗

- What is memory?
- What are pointers and arrays?
- What is `malloc()`?
- Performance discussions of `structs` and functions
- `unions`