

Memory Allocations (in C), Preprocessor, and Branch Prediction

Colin Bretl, Nikolai Maas, Peter Maucher | 24.10.2024

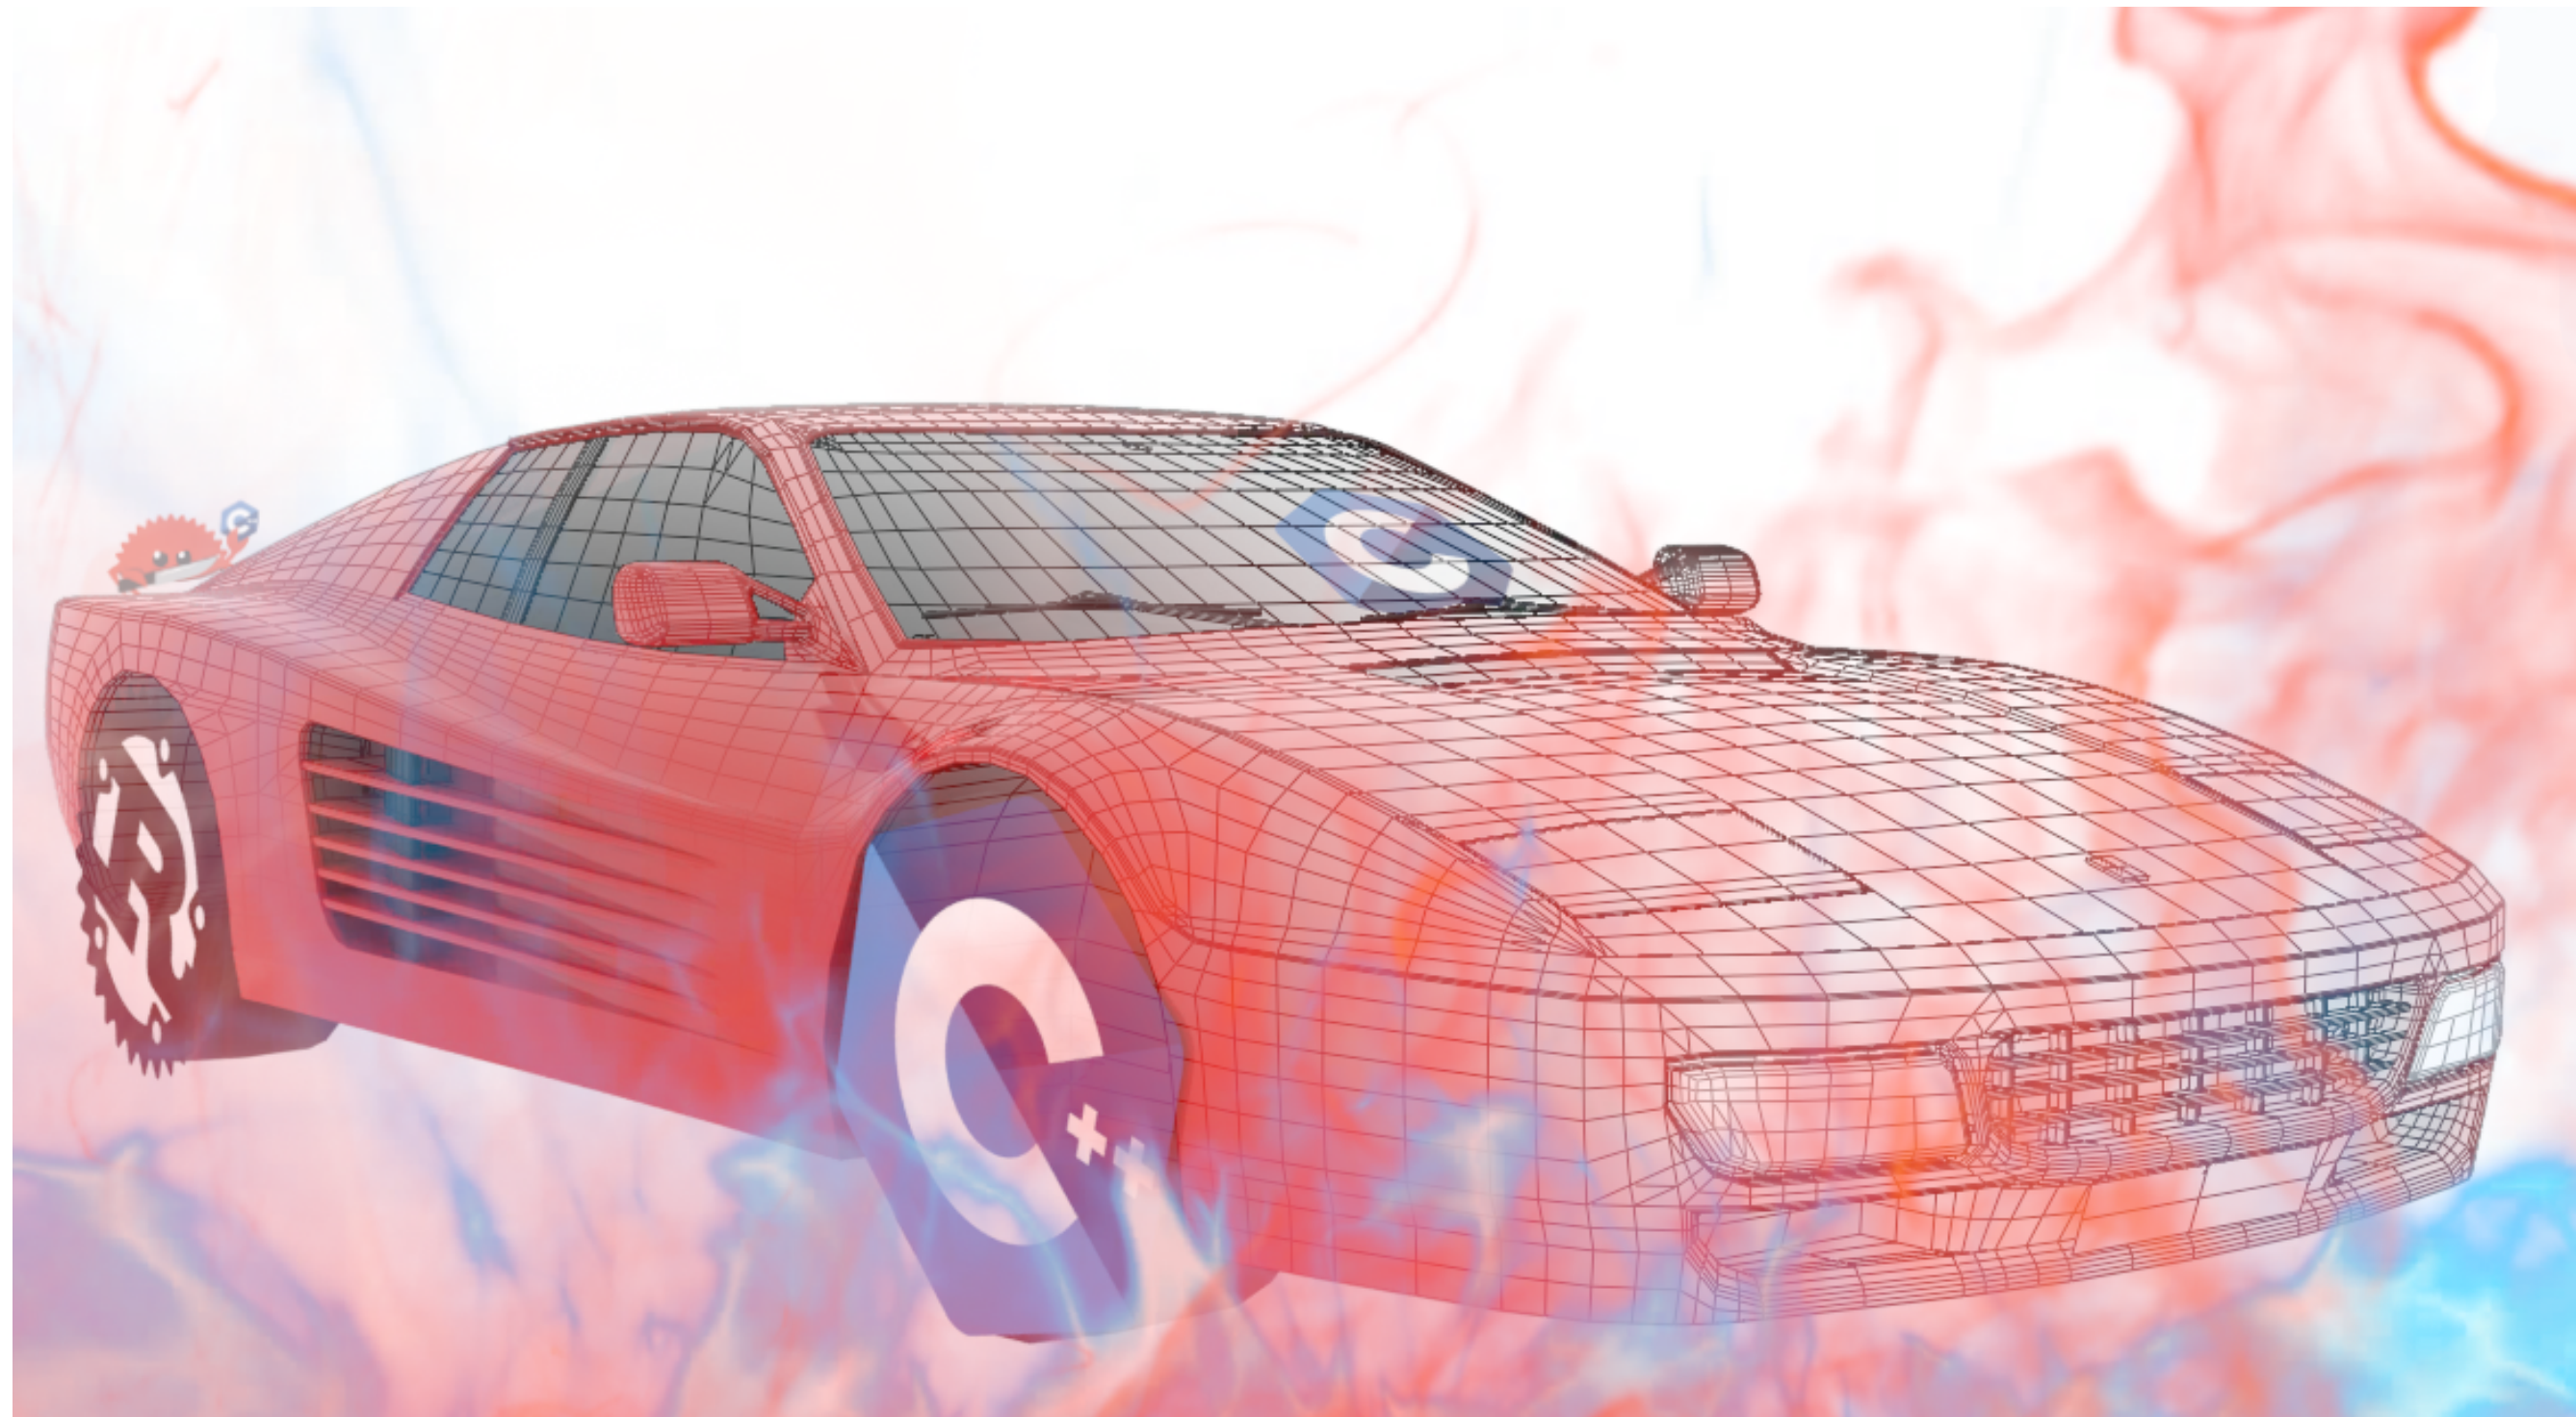


Table of Contents

1. Last Lecture
2. Branch Prediction
3. Allocating Memory
4. Allocation Strategies
5. `union`
6. Recap

Pointing Out the Memory

C RECAP

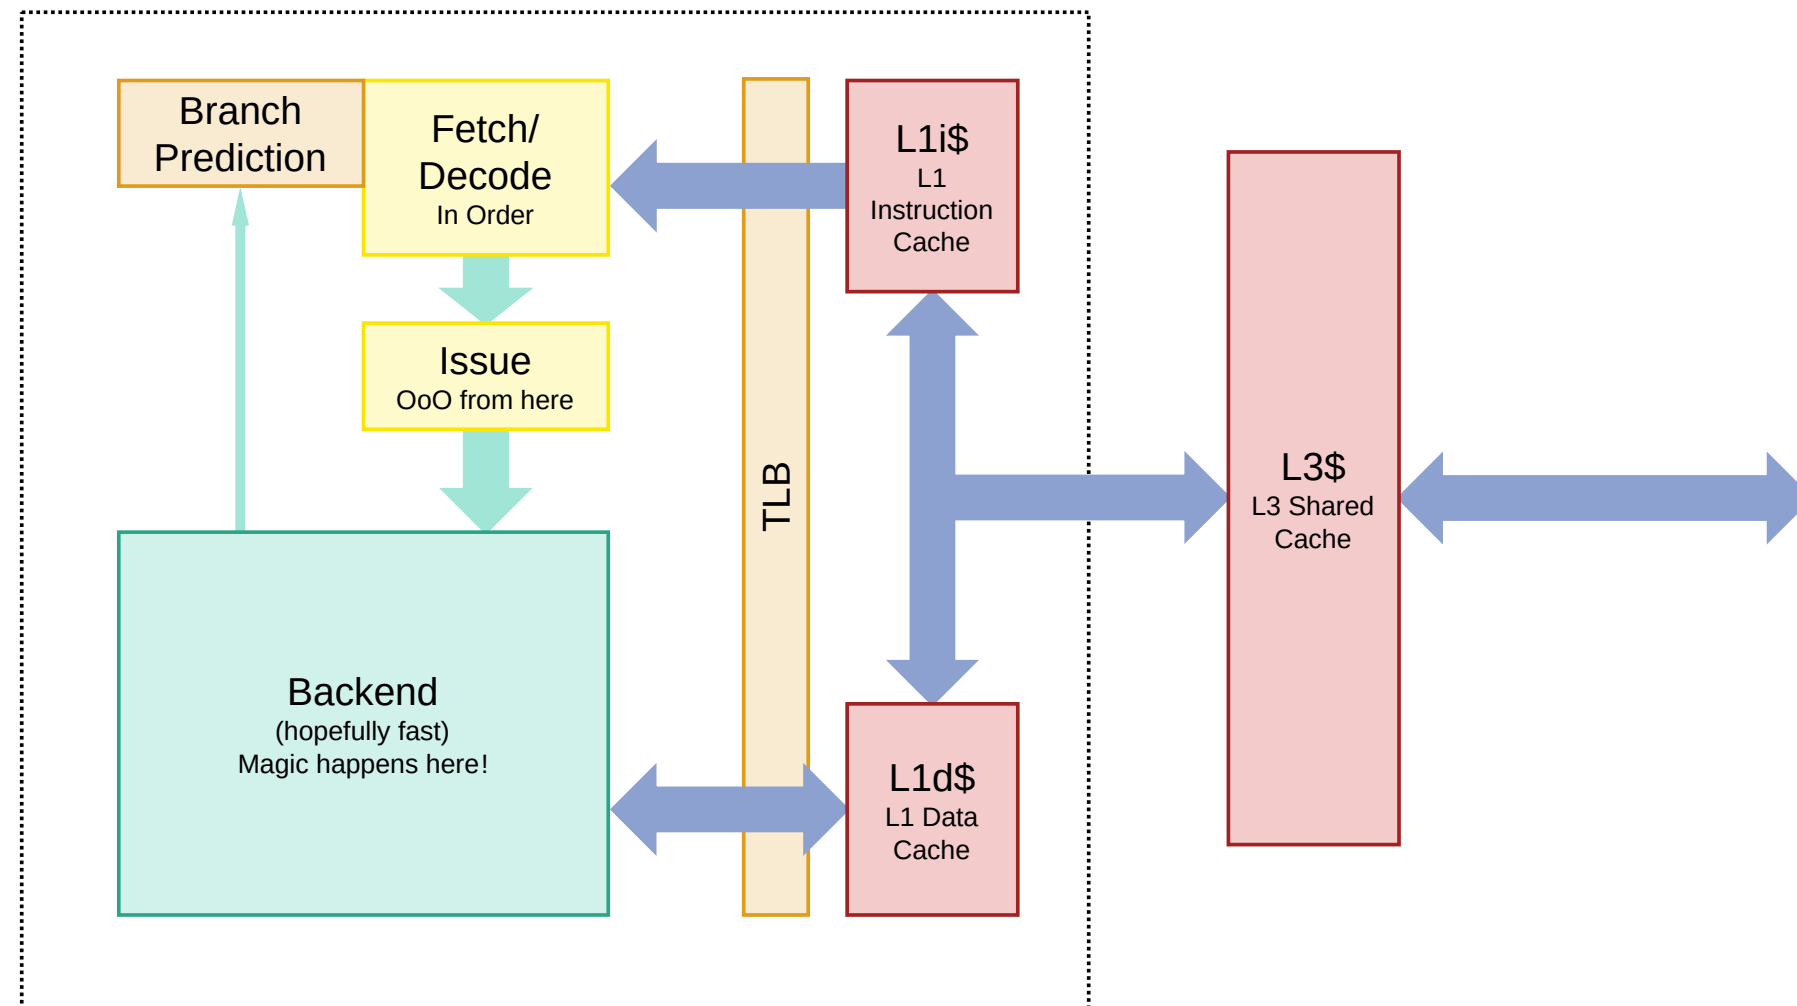
- Types
- Functions
- `struct`

MEMORY

- Pointers
- Caches
- Prefetchers

Branch Prediction

To Branch Or Not to Branch, That Is the Question



Again: μEffCpp

BRANCH PREDICTION

- Decides whether branch is predicted to be taken or not
- Learns from recent occurrences
- Easiest case:
 - Branch upwards: taken
 - Branch downwards: not taken

Why meaningful? Loops mostly repeat

- Second most easy example: Same behavior as last time
- Our model: Combines both
 - Unknown taken if upwards
 - Otherwise last behavior

Branch Prediction

Performance Impact

PREDICTION CORRECT

- CPU knew next instruction(s)
- CPU prefetched said instruction(s)
- CPU (partially) decoded & issued said instruction(s)
- CPU (partially) executed said instruction(s)
- CPU successfully verified the branch prediction
- Branch was basically free (except memory)

PREDICTION INCORRECT

- CPU expected other instructions
- CPU's queues (decode, issue, execute) are blocked by those instructions
- CPU wasted cycles on wrong work
- CPU cleans queues
- CPU fetches correct instructions
- Then: see "Prediction Correct"
- About **15-20 cycles** of wrong work in queues + memory

Preprocessor

Glorified sed `'s/foo/bar/g'`

Preprocessor: Text replacement before actual compiler runs

Main uses

- Copy in text from other files
- Name text constants
- Name text replacement functions
- Conditionally insert/remove code

Preprocessor

This >> # << Is Not a Hashtag!

THE USES OF

- `#include`
- `#define`
- `#undef`
- `#if` / `#if defined` / `#ifdef`
- `#else` / `#elif`
- `#endif`
- `#error`
- `#warning`
- `#pragma`

Preprocessor

Lets Cook `#include <Copy_Pasta.h>`

`#include <FILENAME>`:

- removes the line
- instead inputs the verbatim contents of FILENAME

`#include <FILENAME>` vs

`#include "FILENAME"?`

- `<>`: *Usually* system / library files
- `" "`: *Usually* files from current project

Whats in FILENAME?

- *Usually* used for headers

HEADERS?

Collects data to share code with

- Type definitions: What does this type look like?
- function / variable declarations: This element exists

Access Violation: `#define` `private` `public`

`#define` SOMETHING: replace every (following) occurrence of SOMETHING with something else

SIMPLE MACRO

`#define` SOMETHING OTHERTHING

- SOMETHING now called a macro
- SOMETHING can be (nearly) every character sequence
- Newlines need to be escaped using `\` like this
- *Usually* used for constants
- Example: `#define` PI 4

MACRO "FUNCTION"

`#define` SOMETHING(PARAM) OTHERTHING

- SOMETHING can now use parameters
- parameters are verbatimly replaced
- Example: `#define` MAX ((a)>(b)?(a):(b))
- Example call: MAX(i++, j++)
- Replaced to: ((i++)>(j++)?(i++):(j++))
Attention! Modifies i and j

And `#undef` SOMETHING? makes preprocessor forget SOMETHING

Preprocessor

`#if I_COULD_FLY`

`#if SOMETHING`: Only insert if SOMETHING is `!= 0`

- Preprocessor: No (explicit) parentheses
- End of `#if` directive: `#endif`
- Has `#else` and `#elif`
- Can be nested! (Comments cannot be nested)

```
#if 0
    #if 1
        // this comment dies
    #endif
#else
    #if 1
        // this comment survives
    #endif
#endif
```

`#if defined`

- Often: Do not care about macro value, just existence
- `defined(MACRONAME)`-Operator to the rescue
- Shorthand: `#ifdef`

#warning lecture new, expect mishaps

#error fatal

- Handy if you know you cannot compile

```
#if (__STDC_VERSION__ < 201112L)
    #error this code requires C11 for multithreading
#endif
```

#warning non fatal

- Handy to communicate with programmer

```
#if defined(_MSC_VER)
    #warning microsoft compilers are untested
#endif
```

Not Preprocessor

`#pragma` GCC diagnostic ignored "-Wuninitialized"

`#pragma` is not (really) a preprocessor directive

Allocating Memory

Actually Initializing Our Vectors

STACK

- Local memory of function
- Cannot be **returned** (!)
(Cleaned up on scope exit)
- Passing into called function
ok? Yes
- Fixed size per function
- 🚴 Performance: Extremely **hot**, basically always in L1 Cache ⇒ use often

HEAP

- Global memory, shared
- Can be **return**
- Dynamically sized
- 🚴 Performance: Depends on usage pattern

STATIC MEMORY

- "Forbidden" Category...
- Use with extreme caution
- 🚴 Performance: Similar to heap

Stacks and Scopes

Where Variables Deserve to Die

```
int main() {  
    struct ll_vector last = {0, 2};  
    struct ll_vector mid = {&last, 1};  
    { // explicit scope  
        struct ll_vector first = {&mid, 0};  
    } // first stops living here  
} // last, mid stop living here
```

- Scope: "area" of common lifetime of variables
- Lifetime: time of valid use
- Lifetime: Starts at definition, ends at enclosing scope close
- Scopes (in C): basically everything between { and }
- Scopes: functions, loop/if bodies
- Not a scope: `struct` body

- (CPU) stack: More capable than Algo 1 stack data structure
 - New variable: *pushed* to end (usually low address)
 - After scope exit: *popped* from end
 - (Hence the name)
- Additionally: random access in stack frame (area of function)
- Recursion: One function has multiple stack frames, one for each recursive call

Heaps

Heap, Heap, I'm a Sheep

```
int *iota(int num, int start) {  
    /* create values start, start + 1,  
    start + 2, ..., start + num - 1*/  
    int *values = // help!!!  
}
```

What kind of memory do we want?

- Persists over function boundaries
- Can be "safely" shared across threads
- Type-independent
- Dynamically sized
- "Leak-free"
- Fast allocations

C solution:

```
void *malloc(size_t num_bytes);  
void free(void *pointer);
```

Issues?

- **Alignment:**
`char arr[sizeof(long long)];` and `long long` have same size, but different alignment
 - Solution: Assume worst-case alignment
- **Locality:** No option to allocate, e.g., list in adjacent memory
 - Solution: No real solution
- **Fragmentation:** No option to allocate memory with similar lifetimes in similar memory
 - Solution: fancy heuristics in fast malloc implementations

free ()-styler!

DESIGN IMPLICATIONS

- **malloc ()** needs to go into **free ()**
 - Only information: the pointer
 - We need to track all sizes of allocated memory
 - Store somewhere (e.g., around allocated memory)
- Need to reuse memory: find previously **free ()**d memory
 - Need a lookup structure
 - For fast implementations: acceleration structure
- Need to track whether new memory needs to be requested \Rightarrow OS!
 - Issue: Fragmented size does not count, only contiguous free memory
- Need to actually get memory from the OS...

I want to `brk()` free

`mmap()` men, `mmap()` men, `mmap()` `mmap()` `mmap()` men

GETTING MEMORY FROM THE (LINUX) OS

SBRK()

`sbrk()` sets "program break", the size of the shared data segment

- Can be used for small allocations
- Needs data structure that shows whether lowest data can be freed to shrink "break"
- Big performance issue: Extremely diverse size of allocations (1 - 128 kiB), similar sizes should be adjacent to combat fragmentation and to speed-up lookups.

MMAP()

`mmap()` maps files (and chunks of memory) into address space

- Can be used for all allocations
- Technical limitation: multiples of page size (usually 4 kiB)
- OS needs exact information on class="text-unimportant" `munmap()`
- Setup security features like "Guard Pages"
- Performance issue: demand paging and pre paging

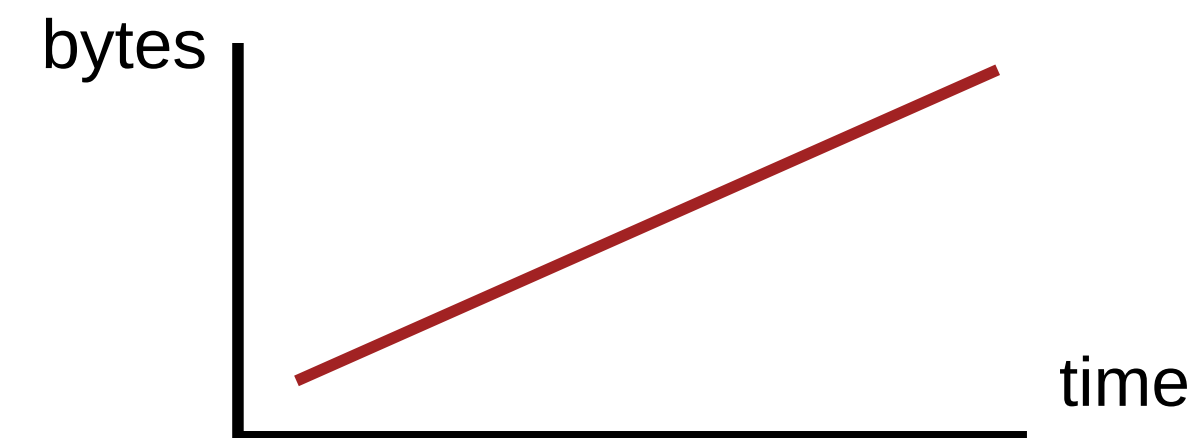
Embrace the Penguin

- Does `malloc ()` fail? No (except in OS assignments). Still check!
- Why not? Demand Paging
- What happens on out of memory? Out of Memory (OoM) killer
- What is targeted by OoM killer? Applications with a lot of memory, e.g., the X-server, the desktop manager, or *your application...*
- What do you do then? End other applications or get more RAM

Allocation Strategies

OS *copy = malloc(pasta): Allocation Patterns

- Ramp
 - Allocate more and more, does not free
 - Example: file system cache
- Peaks
 - Allocate in phases
 - Example: pipelined processes
- Plateau
 - Does not allocate, (only at beginning)
 - Example: long computations



Applications have some typical allocation pattern
adapted from source

Allocation Strategies

Please a Bit - Bitmap Allocator

- Pick a size: e.g., 64 byte (cache line!)
- Reserve a memory range for actual storage
- Reserve a memory space where you store one bit per byte
- Allocate by rounding up size, finding consecutive free block large enough
- Set bits for "allocated"
- On free? Need to find length
- Malloc-style allocator: size lookup somewhere else
- Other, great idea: use only for sizes \leq blocksize, free becomes trivial

PERFORMANCE

- Bit range: extremely hot \Rightarrow great
- size lookup for free() slow
- If single size (or range, e.g., 32-64 bytes): Extremely fast \Rightarrow *Slab Cache*
- Programs often allocate same size often

Free List

- Linked List: [next, length]-segments
- Store in returned memory (no space overhead!)
- On alloc: iterate through list until length \geq requested length
- Remove from list, return segment
- Possibly add new, smaller block to list
- Issue: alignment
- On free? Need to find length
- Simple idea: store before returned pointer (mind the extra space)
- Remember to merge segments if possible

PERFORMANCE

- Implementation advantage: Simple to implement, ok-ish for all allocation sizes
- Performance issue: linked-list
- Expectation: Newly returned memory hot.
Usually wrong, applications defer free()
- Remember? Programs often allocate same size often
- Usually, first try succeeds \Rightarrow actually decent!

Allocation Strategies

My Best Buddy - Buddy Allocator

- Allocatable memory size: $2^n, n \in \mathbb{N}$
- Allocation request: round up to next power of two
- Free-list of all blocks of required size. If empty: (recursively) allocate one power larger, split result and insert other block (= buddy) into free-list
- On free: insert now freed block into free-list. If buddy is also free, merge (recursively) to regain large blocks
- Issues:
 - Internal & external fragmentation
 - Space waste up to $\frac{1}{2} - \varepsilon$

PERFORMANCE

- Possibly replace free-list with other tool.
- Advantage: Simple to implement
- Performance: binary split really fast
- Depending on free-list or replacement: merge slow(ish)
- Not (really) general purpose, but great building block!

Allocation Strategies

Welcome To the Arena

- Remember application phases? Wouldn't it be great to only cleanup once at the end of the stage?
- Idea: bin of memory (with parents)
- Allocate to bin that lives as long as needed, but not longer
- At the end of live:
`deallocate_all(arena *)`
- Here: no real implementation advice, depends on your use case

PERFORMANCE

- Hugely depends on correct usage, but then can be really fast
- No automatic allocation to correct arena possible, needs programmer input
- But does it? Machine learning to the rescue!
 - Learning-based Memory Allocation for C++ Server Workloads
 - Combining Machine Learning and Lifetime-Based Resource Management for Memory Allocation and Beyond
 - Performance Test Selection Using Machine Learning and a Study of Binning Effect in Memory Allocators

Allocation Strategies

Composition is Key

- No allocator great alone
- But together: great building blocks
- General Pattern: Different behavior based on size (and alignment)

Example

- If $\text{size} \geq 2\text{MB} \Rightarrow$ call `mmap()` directly
- If $2\text{MB} > \text{size} \geq 128\text{kB} \Rightarrow$ Use page-aligned free-list backed by `mmap()`
- All below: backed by buddy allocator, all buddy-managed memory initially 2MB aligned
- If $128\text{kB} > \text{size} \geq 4\text{kB} \Rightarrow$ Use different page free-list
- If $4\text{kB} > \text{size} > 128\text{ B} \Rightarrow$ Use 64 B-aligned free-list
- If $128\text{ B} \geq \text{size} > 64\text{ B} \Rightarrow$ Use 128 B slab cache
- If $64\text{ B} \geq \text{size} > 48\text{ B} \Rightarrow$ Use 64 B slab cache
- If $48\text{ B} \geq \text{size} > 32\text{ B} \Rightarrow$ Use 48 B slab cache
- If $32\text{ B} \geq \text{size} > 16\text{ B} \Rightarrow$ Use 32 B slab cache
- If $16\text{ B} \geq \text{size} \Rightarrow$ Use 16 B slab cache

United We Stand, Devided We Undefined Behavior (UB)

```
union foo {  
    long long lvalue;  
    double dvalue;  
};
```

- Either store lvalue or dvalue
- Can NOT be used to convert (at least not savely)
- Only read from object last written to!
- Basically only useful for saving memory
- And if we want to convert?
- use `memcpy()`
- Performance🚴: Pretty quick, but use with caution due to UB

Recap

What did we learn today?

- What is the branch predictor?
- What is the preprocessor?
- `malloc` and `free()`
- Where and how do we allocate memory?
- What is a `union`?

And next lecture? Not a lecture, first practical session

- Practical introduction
- Repository creation
- Exercise introduction