

Effizientes Programmieren in C, C++ und Rust

Exceptional C++

Colin Bretl, Nikolai Maas, Peter Maucher | 24.10.2024



Table of Contents

1. Failing In Place
2. Designing Our Way Out
3. Keyword Matching
4. `std::exception`
5. `whoami`
6. `dynamic_cast`
7. `noexcept`
8. Performance
9. Error handling performance
10. The Final Result
11. Templates
12. Generic Types
13. Recap

Failing In Place

Add a Way Out

- Constructors can fail
- Constructors are *not* functions
 - Constructors do *not* return
 - Constructors have initializer lists
 - Constructors run in-place
 - Constructors ignore `const`
 - Constructors cannot *fail partially*
- Where failure in construction?
 - In subconstructors
 - In own code
- On failure in construction: roll-back, abort, and report
 - Undo own code (if we got so far)
 - Call destructors of already called constructors (but *only* those)
 - Abort and report failure

Designing Our Way Out

The Road to Error Handling is Paved With Good Options

- **Hard qualifications** of mechanism:
 - Separate in normal and exceptional state
 - Find caller that handles report \Rightarrow stack unwinding
 - In exceptional state: destroy all destructors on path to handler
 - Report arbitrary error information
- **Soft qualifications** of mechanism:
 - Fast for normal (non-failure) path
 - Clear separation between production code and error handling
 - Nice syntax

Given that our mechanisms deals with the *exceptional* code path, let's call it **exception**

Keyword Matching

finally Something that C++ Does Not Support

- Entering the exceptional state: **throw**
 - Enters stack unwinding procedure
 - Can have one object of metadata
 - Any metadata: `throw -1;` is absolutely valid
 - Later today: `std::exception`
- Checking for exceptions:

```
try { throw Ball {}; } catch (Ball&) {}
```

 - Checks exception (almost always by reference!)
 - Additionally checks for exception type (RunTime Type Information, RTTI)
 - Multiple **catch** blocks possible
 - Last **catch** (...) can rethrow without parameters
- Not in the language: **finally**
 - Use RAI!
- Also not in the language:

```
catch ... else
```

 - Use longer **try** { /* */ } clause

`std::exception`

The exception base class to be used

- `std::exception`
- The base class for all your exception needs
- Probably: Select one of the subclasses (e.g. `std::runtime_error`)
- Use the standard library! ↗
- See the part with "do not exit with an exception"?
- Throwing exceptions while an exception is thrown is bad
- ***Do Not Throw In A Destructor!***
- Chaos!

RTTI

- Sometimes: object with unknown runtime type
 - Reference or Pointer only
 - "Polymorphic Type"
- Query static type: `typeid`
- Result: `std::type_info`
 - Comparison: Are these both for the same type?
 - `before()` allows sorting of types, great
 - `hash_code()` allows hashing of types, great
 - `name()` some name that correlates with the types' name

```
#include <iostream>

struct base {};
struct derived : base {};

struct base2 {virtual ~base2() = default;};
struct derived2 : base2 {};

int main() {
    derived d;
    derived2 d2;
    std::cout << sizeof(d)
              << " " << sizeof(d2) << std::endl;
    base *b = &d;
    base2 *b2 = &d2;
    std::cout << typeid(*b).name()
              << " " << typeid(*b2).name() << std::endl;
}

g++ typeid.cc && ./a.out
1 8
4base 8derived2
```


Verifying runtime types

- Often: child classes have more info
 - Base class: Drawable, children: Rectangle, Circle
- When drawing: Need extra info
- Use `typeid()` for dynamic query? For selection only
- We need to cast

- Use `dynamic_cast`
- Two options:
 - With pointers, returns `nullptr` if wrong type
 - With references, throws exception

```
derived2 *d2p = dynamic_cast<derived2*>(b2);  
derived2 &d2r = dynamic_cast<derived2&>(*b2);  
base2 bc{};  
derived2 &d2rf = dynamic_cast<derived2&>(bc);
```

```
terminate called after throwing an instance  
of 'std::bad_cast' what(): std::bad_cast  
[1] 883256 IOT instruction (core dumped) ./a.out
```

Promises, Promises

- Realistically: Most functions do not throw
 - Pure Math
 - Unrecoverable Errors (out of memory, CPU crashed, ...)
 - Someone used wrong error handling (`std::filesystem`)
- In this case: Promise not to throw
- Enables optimizations and better code understanding
- `int add1(int a) noexcept;`
- As usual in C++: wrong default `_(ツ)_/`

- *Historical aside: checked exceptions*
- *In C++98: Possibility to explicitly list all throwable exceptions*
- If unexpected exception, call `std::unexpected`, which terminates
- Idea: Simplify arguing about code, increase performance 🚴
- Reality: Not really used
 - Unable to modify code locally
 - Not enforced by compiler, but at runtime
 - Only used with empty list (can throw no exception)
 - Nowadays: **checked exceptions considered harmful**

Are Exceptions Slow?

Implementing Exceptions Efficiently

- Exception implementation guidelines: 🚲
 - Rule No 1: Optimize the normal, error-free path
 - Rule No 2: The error path can be very slow
 - Rule No 3: See rule 1
- Rough idea: Emit code as if no exception could ever occur
- On throw: Jump to completely different part of code ("shadow stack", Peters name)
- Per function: Store the start of various destructors, catch clauses
- Parse return stack, pc to jump to correct next function handler on shadow stack
- Sometimes: requires use of additional base pointer
- Overhead of exceptions (when not thrown): About one instruction per function

Error handling performance

It's Just Bad

Error handling is never free!

- Optional: Consumes extra stack space, destroys register allocation (calling convention)
- Error Codes: destroy return value, require extra pointer and stack
- Confusing normal and exceptional path, basically universal

The Final Result

Expect the expected

- What we really want:
 - Type that contains *either* the expected result *or* an error
 - Basically an improved `std::optional`
 - Exception-free code as long as we're dealing with the expected or with the error explicitly
 - If we cannot handle the unexpected locally \Rightarrow throw the unexpected
- See: `std::expected`
- *Basically an improved `std::optional`*
- *C++23 only :/*

```
#include <expected>
#include <string>
#include <iostream>

std::expected<int, std::string> parse_char(char c) {
    using namespace std::string_literals;
    if (c < '0' || c > '9')
        return std::unexpected{"character is not a number: "s + c};
    return c - '0';
}

auto inner() {
    auto first = parse_char('5');
    auto second = parse_char('a');
    if (!first.has_value()) return 0;
    return first.value() + second.value();
}

int main() {
    try {
        inner();
    } catch (std::bad_expected_access<std::string> &err) {
        std::cout << err.error() << std::endl;
    }
}
```

Templates

And Now For Something Completely Different

- Often: Write code that is type-agnostic
- e.g., sorting
 - Need sizeof (for pointer arithmetic)
 - Comparison
- C approach: **copy-paste code**
- C++ approach: **let compiler** replace types
 - Create class or function *template*
 - Compiler creates new element each time using template
 - Called *template instantiation*
 - Rust name: *Monomorphization*

```
#include <iostream>
#include <utility>
#include <vector>

template<typename T>
auto bubble_sort(std::vector<T> input) {
    if (input.size() < 2) return input;
    for (auto i = input.size(); i > 1; i -= 1) {
        for (auto j = 0; j < i - 1; ++j) {
            if (input[j] > input[j + 1])
                std::swap(input[j], input[j + 1]);
        }
    }
    return input;
}

int main() {
    auto vec1 = std::vector
        {5, 4, 3, 2, 1, 6, 7, 8, 9};
    auto vec2 = std::vector
        {5.5, 4.5, 3.5, 2.5, 1.5, 6.5, 7.5, 8.5, 9.5};
    std::cout << (typeid(vec1) == typeid(vec2)) << std::endl;
    vec1 = bubble_sort(vec1);
    vec2 = bubble_sort(vec2);
    for(auto e1 : vec1) std::cout << e1 << ", ";
    std::cout << std::endl;
    for(auto e1 : vec2) std::cout << e1 << ", ";
    std::cout << std::endl;
}
```


Generic Types

Writing Libraries, Properly

- Often: Write code that is type-agnostic
- Extremely often: data structures (i.e., containers)
 - Contain more objects, handle them
 - list, vector, ...: completely data agnostic
 - tree, hashmap: require some data functions (concepts, eventually)
- Even specialized code (Quadtree, bitsets, ...) can decide on data width
- Almost everything can be a template!
- Decision guide: **template if actually required, else keep refactorable**

```
#include <iostream>
#include <memory>

template<typename T>
struct list {
    std::unique_ptr<list<T>> next;
    T data;
    void append(T data) {
        auto ptr = this;
        while(ptr->next)
            ptr = ptr->next.get();
        ptr->next = std::make_unique<list<T>>(std::move(data));
    }
    list() = default;
    list(T data) : list {nullptr, data} {}
private:
    list(std::nullptr_t, T data) : next {nullptr}, data{data} {}
};

int main() {
    list<int> l{0};
    for (auto i : {1, 2, 3, 4})
        l.append(i);
    auto ptr = &l;
    while(true) {
        std::cout << ptr->data << ", ";
        if (ptr->next)
            ptr = ptr->next.get();
        else break;
    }
    std::cout << std::endl;
}
```

Recap

What we did today

- Today:
 - Exceptions
 - Exception performance 🚴
 - Exception alternatives
 - Template introduction
- Next Week:
 - Template in depth