


# Effizientes Programmieren in C, C++ und Rust

## Rust - Introduction and Basics

Colin Bretl, Nikolai Maas, Peter Maucher | 24.10.2024



# Recap

1. What did you like about C++?
2. What did you hate about C++?

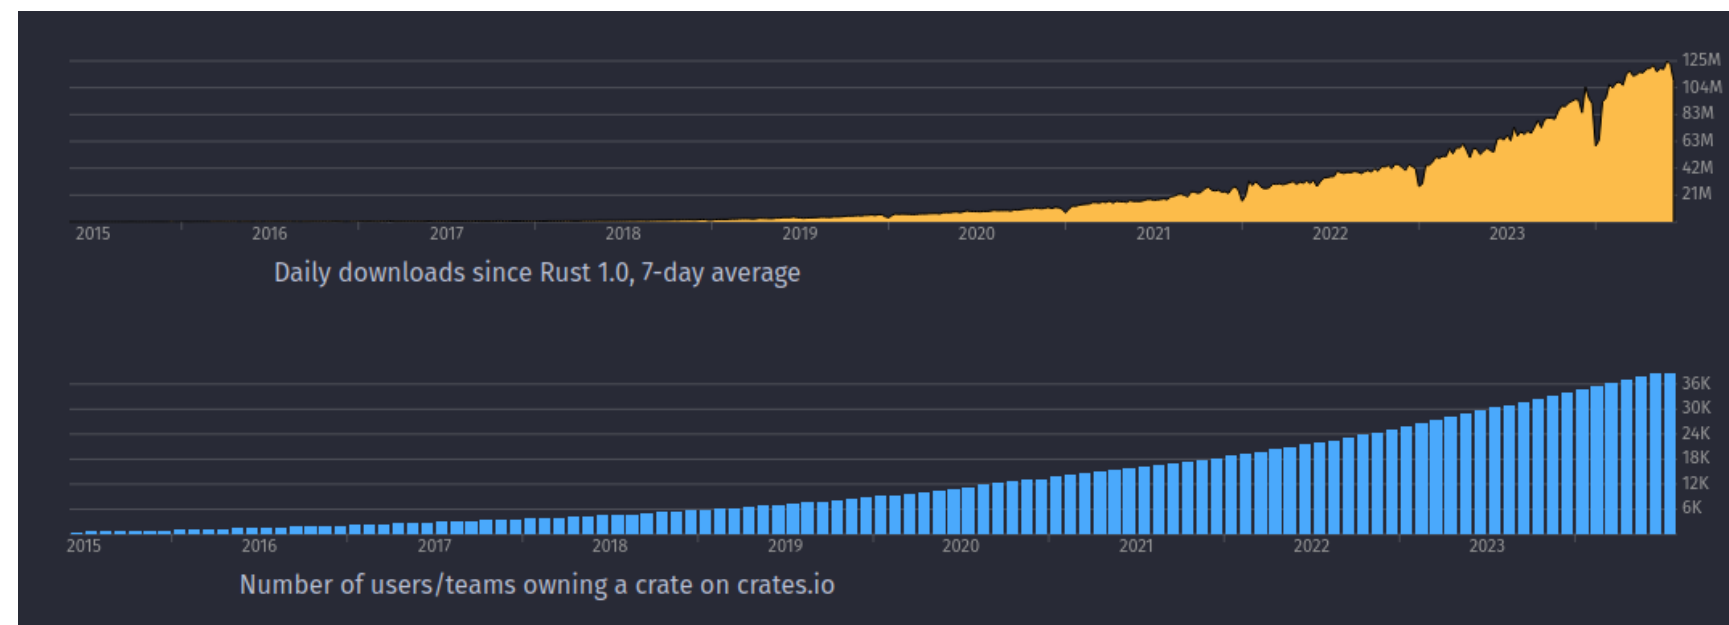
# Table of Contents

1. A New, Hot Systems Programming Language
2. Safety vs Performance
3. Rust: Standing on the Shoulders of Giants
4. Getting Started
5. Structs
6. Enums and Pattern Matching
7. Pattern Matching versus Dynamic Dispatch
8. Nullability, or: Maybe Just Option
9. Error Handling in Rust
10. Modules
11. Macros
12. More about the Rust Toolchain
13. Summary

# A New, Hot Systems Programming Language

- Probably the most hyped programming language at the moment
- Fast growing
- Serious competition for C++?
- RIIR: rewrite it in Rust
  - *almost a technical term by now*

style="width: 85%" /> style="width: 85%" />



style="width: 50%;"/>

<https://lib.rs/stats>

# What is behind the hype?

## The Official Website

"A language empowering everyone to build reliable and efficient software."

"**blazingly fast** and memory-efficient"

"memory-safety and thread-safety"

"top-notch tooling"

"friendly compiler with **useful error messages**"


## Goals

- Similar to C++: high-level, efficient systems programming
  - ⇒ Strong abstraction capabilities, same performance as C
- Why can't we just use C++?
  - C++ is full of **footguns** (*memory safety/undefined behavior*)
  - C++ is burdened with a giant legacy of backwards-compatibility




# Safety vs Performance

- C++: we pay for high efficiency with undefined behavior (UB) ↗
- **Traditional view:** either safety or low level control and efficiency
- **Rust:** Why don't we have both?  
*(with some caveats)*



## Achieving Safety and Performance

- Separation into **safe** and **unsafe** Rust  
⇒ All the UB lives exclusively in unsafe
- Unsafe code is wrapped in a safe API
- Powerful type system (borrow checker!) ensures the API can't be misused

"Rust locks all the footguns in a safe – and un-safe is the key" 

⇒ **Principle:** build safe abstractions upon an unsafe fundament

# Rust: Standing on the Shoulders of Giants

- Development started (officially) in 2009 at Mozilla
  - *Hearsay: the developers were too annoyed with the problems caused by C++ in Firefox*
- Rapid evolution, stable 1.0 release in 2015

## So what is Rust?

- "C++ done right"
  - ⇒ Zero-cost abstractions, monomorphization, value/reference-semantics, ...
- Strong functional influences
  - ⇒ Pattern matching, immutability (default), ...
- Ownership and the borrow checker!



- **Claim 1:** Rust is both object oriented and functional
- **Claim 2:** Rust is neither object oriented nor functional

# Getting Started

- Unified build system and dependency management: cargo

## Rust Hello World

```
ferris@rustbootcamp: cargo new hello-world  
ferris@rustbootcamp: cd hello_world && tree
```

```
├── Cargo.toml  
└── src  
    └── main.rs
```

### main.rs

```
fn main() {  
    println!("Hello, world!");  
}
```

## Build and Run

```
ferris@rustbootcamp: cargo build --release  
    Compiling hello_world v0.1.0 (/home/.../hello_world)  
    Finished release [optimized] target(s) in 0.11s  
ferris@rustbootcamp: ./target/release/hello_world  
Hello, world!
```



# Getting Started

- Recommended reading: the book ([doc.rust-lang.org/book/](https://doc.rust-lang.org/book/))
- Control flow syntax similar to C++
- **Expression oriented**: every statement returns a value (e.g. if-else)
- Type inference for local variables
- `double` → `f64`, `int` → `i32`, `size_t` → `usize`
- **Immutable by default**: `mut` required to mark variable as mutable
- Additional primitive types:
  - Tuples
  - Arrays (fixed size)
  - Slices (variable size, non-owned)

```
fn max_of_slice(vals: &[f64]) -> f64 {  
    let mut max = f64::MIN;  
    for &val in vals {  
        max = if val > max { val } else { max };  
    }  
    max // <- implicit return  
}
```

```
let array: [i32; 10] = [0; 10];  
let slice: &[i32] = &array; // get slice of array  
  
let result: (i32, i32) = min_max_of_slice(slice);  
let (min, max) = result; // destructuring  
println!("input={}, min={}, max={}",  
        array, min, max);
```

# Getting Started

## References and Move Semantics

- Borrowing: use `&/&mut` to get a shared/unique reference
- The borrow checker ensures only valid references are used!
- Default assignment **moves** values
- Can't be used afterwards, invalidates references
- Exception: Some types are copied instead (e.g. numbers, plain structs)

⇒ More details later

```
let mut my_vec = Vec::<i32>::new(); // *
my_vec.push(42);
let answer = &my_vec[0]; // reference first element
do_something_else(my_vec);
// use after free?!
println!("the answer is {}", answer);
```

```
error[E0505]: cannot move out of `my_vec` because it is borrowed
--> src/main.rs:24:23

21 |     let mut my_vec = Vec::<i32>::new(); // <- turbofish!
    |         ----- binding `my_vec` declared here
22 |     my_vec.push(42);
23 |     let answer = &my_vec[0]; // reference to first element
    |                   ----- borrow of `my_vec` occurs here
24 |     do_something_else(my_vec);
    |                       ^^^^^^ move out of `my_vec` occurs here
25 |     println!("the answer is {}", answer)
    |                               ----- borrow later used here
```

*\* this syntax is called a turbofish  
(and avoids the syntactic ambiguity of C++ templates)*

# Structs

- Basic purpose analogous to C++

## Syntax

- **Self** (upper case) is the type of the current struct
- **self** parameter indicates methods
- Member access only with explicit **self**
  - *Rust generally prefers explicitness*
- Initialization via **Self** { ... }

```
struct Vector2D { x: f64, y: f64 }

impl Vector2D {
    // "static" function for construction
    pub fn new(x: f64, y: f64) -> Self {
        Self { x, y } // initialization
        // ^- shorthand for Vector2D { x: x, y: y }
    }

    // method (shared reference)
    pub fn length(&self) -> f64 {
        (self.x * self.x + self.y * self.y).sqrt()
    }

    // method (unique reference)
    pub fn set(&mut self, x: f64, y: f64) {
        self.x = x;
        self.y = y;
    }
}
```

# Structs are Not Classes

- No inheritance, no unique address
- No metadata or vtable-pointer (POD)
- Only implicit "constructor"
  - Structs always fully initialized (C++: partial initialization observable)
  - Use functions to define different ways of construction
- Methods are syntactic sugar for functions
  - *Also: atomic dereferencing (no `->` operator)*
  - *Dynamic dispatch only via traits ↗*
- Primitive types also support methods

```
struct Vector2D { x: f64, y: f64 }

impl Vector2D {
    pub fn new(x: f64, y: f64) -> Self {
        Self { x, y }
        // ^- shorthand for Vector2D { x: x, y: y }
    }

    pub fn length(&self) -> f64 {
        (self.x * self.x + self.y * self.y).sqrt()
    }
}
```

```
let v = Vector2D::new(1.0, 2.0);
assert_eq!(
    v.length(), Vector2D::length(&v)
);
```

```
assert_eq!(1.max(2), i32::max(1, 2));
```

⇒ **No difference in behavior & efficiency**  
between primitives and structs!

# Enums and Pattern Matching

- Usually called **sum types** or **tagged unions**  
 $\Rightarrow$  functional programming!
- Represents one of multiple variants
  - *set theoretic view: member of  $A + B + C$*   
*for variant sets  $A, B, C$*
- Consist of associated data and **tag**
- Support methods etc. in the same way as structs
- Used via **pattern matching** (**match**)
  - "switch on steroids"
- C++: no language support,  
`std::variant` as poor substitute

```
enum Message {
    Quit, // unit variant
    Move { x: i32, y: i32 }, // struct variant
    ChangeColor(u8, u8, u8), // tuple variant
}

impl Message {
    fn handle(&self) {
        match self {
            Message::Quit => process::exit(0),
            Message::Move { x, y }
                => println!("move to {x}, {y}"),
            Message::ChangeColor(r, g, b)
                => println!("#{r:02X}{g:02X}{b:02X}"),
        }
    }
}
```



# Enums and Pattern Matching

- Pattern matching is powerful and versatile
    - Nested patterns, arbitrary conditions
    - Compiler ensures no case is missing
  - Example: match slice and contained enum variants at once
  - Also usable via destructuring, **if let** and **let ... else**
- ⇒ But is the resulting code **efficient**?

```
fn handle_multiple(message_list: &[Message]) {  
    match message_list {  
        [msg] => msg.handle(),  
        [Message::Quit, ..] => process::exit(0),  
        [head, tail @ ..] => {  
            head.handle();  
            handle_multiple(tail);  
        }  
        _ => { }, // catch-all  
    }  
}
```

```
fn handle_move(msg: Message) {  
    let Move{ x, y } = msg else {  
        return;  
    };  
    // do something with the move  
}
```

# Enums and Pattern Matching

- What are the performance implications? 🚴

## Data Layout


- **enum** = data + tag
- Tag must be aligned!  $\Rightarrow$  more than one byte
- Data size is defined by the **largest variant**  
 $\Rightarrow$  Seldomly used large variants problematic  
– move to heap?
- **Niche optimization**: compiler uses invalid states of the data to eliminate the tag

```
enum Niche { A(bool), B, C }

println!("{}", mem::size_of::<Niche>());
// 1
```

```
enum Message {
    Quit,
    Move { x: i32, y: i32 },
    ChangeColor(u8, u8, u8),
}

println!("{}", mem::size_of::<Message>());
// 12
```



# Enums and Pattern Matching

## Code Generation

- Naive approach: check every condition one by one (**if else**-chain)
  - *Linear in number of match arms*
  - *Branch mispredictions?*

## Some test results\*

- Simple mapping of variant to constant creates a **lookup table**
- Compiler tries to generate a **jump table** starting at 4 arms
- For **if else**-chains, the compiler tends to go top to bottom in order

```
match message_list {  
  [msg] => msg.handle(),  
  [Message::Quit, ..] => std::process::exit(0),  
  [head, tail @ ..] => {  
    head.handle();  
    handle_multiple(tail);  
  }  
  _ => { }, // catch-all  
}
```

- In more complicated cases, jump tables are not possible  
⇒ For **really hot** matches, check generated assembly and try different code layouts

\* *no reproducability guarantee*

# Aside – Lookup Tables and Jump Tables

- First case: map variant to value

```
enum Variants { A, B, C, D }
```

```
match var {  
  Variants::A => -1,  
  Variants::B => 42,  
  Variants::C => 1337,  
  Variants::D => 13,  
}
```

## A Lookup Table

```
# extract tag  
movsx    rax, dil  
  
# load base address  
lea      rcx, [rip + .lookup_table]  
  
# read table entry  
mov      eax, dword ptr [rcx + 4*rax]  
  
.lookup_table:  
    .long 4294967295  
    .long 42  
    .long 1337  
    .long 13
```

- Effectively compiled to an array access
- Array is part of binary
- Requires no branch

# Aside – Lookup Tables and Jump Tables

- Second case: do something more complicated for each variant (e.g. call a function)

```
enum Variants { A, B, C, D }
```

```
match var {
  Variants::A => callA(),
  Variants::B => callB(),
  Variants::C => callC(),
  Variants::D => callD(),
}
```

- **Note:** This is actually not optimal
- Could we save one indirection?

## A Jump Table

```
movzx    eax, dil    # get jump address via table
lea      rcx, [rip + .address_table]
movsxd   rax, dword ptr [rcx + 4*rax]
add      rax, rcx     # address is relative (ASLR)
jmp      rax          # jump to according address
```

```
.address_table:
    .long    .jmp_a-.address_table
    .long    .jmp_b-.address_table
    .long    .jmp_c-.address_table
    .long    .jmp_d-.address_table

.jmp_a:    jmp    callA
.jmp_b:    jmp    callB
.jmp_c:    jmp    callC
.jmp_d:    jmp    callD
```

- Get address of code for current variant from lookup table
- Jump directly to address



# Pattern Matching versus Dynamic Dispatch

- **Scenario:** data with multiple variants
- OOP solution: inheritance
  - ⇒ one subclass per variant, polymorphic methods

```
struct Shape {
    virtual double area() = 0;
};

struct Rectangle: public Shape {
    double x, y;
    double area() override { return x * y; };
};

struct Square: public Shape {
    double width;
    double area() override { return width * width; };
};

struct Circle: public Shape {
    double radius;
    double area() override { M_PI * radius * radius; };
};
```

- FP solution: pattern matching
  - ⇒ represent data as enum variants

```
pub enum Shape {
    Rectangle(f64, f64),
    Square(f64),
    Circle(f64),
}

impl Shape {
    fn area(&self) -> f64 {
        match self {
            Shape::Rectangle(x, y) => x * y,
            Shape::Square(x) => x * x,
            Shape::Circle(r) => PI * r * r,
        }
    }
}
```

# Pattern Matching versus Dynamic Dispatch

## Pattern Matching

- Method is defined in a single place (static function!)
  - ⇒ Can add **method** without modifying existing code
- Flat data layout

```
impl Shape {
    fn area(&self) -> f64 {
        match self {
            Shape::Rectangle(x, y) => x * y,
            Shape::Square(x) => x * x,
            Shape::Circle(r) => PI * r * r,
        }
    }
}
```

## Dynamic Dispatch

- Variant is defined in a single place
  - ⇒ Can add **variant** without modifying existing code
- Requires indirection and a vtable (e.g. via reference or `std::unique_ptr`)

```
struct Rectangle: public Shape {
    double x, y;
    double area() { return x * y; };
};

struct Square: public Shape {
    double width;
    double area() { return width * width; };
};

struct Circle: public Shape {
    double radius;
    double area() { M_PI * radius * radius; };
};
```

# Pattern Matching versus Dynamic Dispatch

## Benchmark (1k)

- Sum area of 1.000 random shapes
- Input: `&[Shape]` versus `std::vector<std::unique_ptr<Shape>>&`
- Result: ~780ns versus ~4800ns

## Reasons?

- Worse data layout
  - *Though `&[Box<Shape>]` still achieves ~900ns*
- Additional vtable indirection
- **Inlining!** (*without inlining: ~4500 ns*)
- ... what about larger data?
  - *1k fits in L1 cache  $\Rightarrow$  memory access cheap*

```
impl Shape {
    fn area(&self) -> f64 {
        match self {
            Shape::Rectangle(x, y) => x * y,
            Shape::Square(x) => x * x,
            Shape::Circle(r) => PI * r * r,
        }
    }
}
```

```
struct Rectangle: public Shape {
    double x, y;
    double area() { return x * y; };
};

struct Square: public Shape {
    double width;
    double area() { return width * width; };
};

struct Circle: public Shape {
    double radius;
    double area() { M_PI * radius * radius; };
};
```

# Pattern Matching versus Dynamic Dispatch

## Benchmark (1000k)

- Sum area of 1 million random shapes
- Only ~25% of throughput for 1k elements, factor 1.25 faster than dynamic dispatch

⇒ Dominated by memory accesses

```
impl Shape {  
    fn area(&self) -> f64 {  
        match self {  
            Shape::Rectangle(x, y) => x * y,  
            Shape::Square(x) => x * x,  
            Shape::Circle(r) => PI * r * r,  
        }  
    }  
}
```

## Benchmark (1000k, randomized)

- Shuffle resulting list randomly
- More than factor 3 faster than dynamic dispatch

⇒ Allocations have random order,  
thus the indirection hurts

```
struct Rectangle: public Shape {  
    double x, y;  
    double area() { return x * y; };  
};  
  
struct Square: public Shape {  
    double width;  
    double area() { return width * width; };  
};  
  
struct Circle: public Shape {  
    double radius;  
    double area() { M_PI * radius * radius; };  
};
```

# Pattern Matching versus Dynamic Dispatch

- Pattern matching is usually more efficient
- Even better, if possible: avoid both

## Software Design Aspects

- Enums are **closed** to new variants and **open** to new operations

⇒ reversed to OO solution

*(see: visitor pattern, expression problem)*

- Arguably, simpler and more maintainable in many cases
- **Impossible** if external code needs to add variants

## Efficiency Aspects

- Flat data layout, while dynamic dispatch requires indirection
- Jump table comparable to vtable
- Dynamic dispatch typically **not inlined!**
- However: large enum variants and long **if else**-chains are performance pitfalls



# Nullability, or: Maybe Just Option

- There is no `null` in (safe) Rust
- What to do instead?  $\Rightarrow$  Enums, of course!
- Also known as: `Optional`, the Maybe monad

## Correctness

- Nullability encoded in type system
- Compiler enforces checks where necessary

```
enum Option<T> {
    Some(T),
    None,
}
```

```
fn twice(val: Option<u32>) -> u32 {
    val + val
}
```

```
error[E0369]: cannot add `Option<u32>` to `Option<u32>`
--> src/main.rs:11:9
11 |     val + val
    |     ^     --- Option<u32>
    |     |
    |     Option<u32>
```

## Efficiency

- Niche optimization: `Option<&T>` is represented as nullable pointer
- Might reduce null checks compared to `nullptr` (depends on usage)

Better:

```
match val {
    Some(v) => v + v
    None => None,
}
```

or

```
val.map(|v| v + v)
```

# Error Handling in Rust

- By the way: we can also replace **exceptions** with enums (mostly)
- Differentiate two kinds: **bugs/unrecoverable** errors and **recoverable** errors

## Panic!

- Mechanism for unrecoverable errors
- Implementation similar to exceptions (stack unwinding)
- **Not** supposed to be caught!
- Example: bounds checking for arrays  
⇒ Panic instead of buffer overflow

```
panic!("Oh no! Guess we abort the program");
```

```
enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

## Result

- Enum representing (potentially) recoverable errors
- **match** instead of **try ... catch**
- **?**-operator for simple error propagation

```
fn execute_fallible_operations(  
    ) -> Result<Data, Error> {  
    let result: Data = op_that_might_fail()?;  
    // ....  
    Ok(result)  
}
```

# Error Handling in Rust

## Result vs Exception

### Advantages 👍

- Encodes failure conditions in the type system
- Programmer is forced to handle errors
- More composable than checked exceptions
- Error path is much faster than stack unwinding

⇒ Very useful for **high reliability** software

⇒ Can be annoying for prototyping/scripting

### Disadvantages 👎

- Encodes failure conditions in the type system
- Programmer is forced to handle errors
  - *Note: you can "ignore" errors with `unwrap`*  
⇒ *panics instead*
- Happy path incurs an additional (easily predictable!) branch

- Modules are a tool for organizing your source code
  - To split code into multiple files, you need modules
  - Allow to trivially find definitions and determine whether imports are used
- Definition via `mod`, usage via `use` ( $\rightarrow$  *next slide*)
- Visibility management via `pub` and `pub(crate)` (*crate = only within the library*)
- Default "private" visibility is within module

## Definitions are Authoritative!

- Source files without corresponding module definition are ignored
  - $\Rightarrow$  Each new file **first** requires a new module definition!
- Different to most other languages, can be confusing at the beginning

## Most Important Rules

- Modules are a **tree** with root either in `main.rs` (binary) or `lib.rs` (library)
  - The root is referenced via **`crate::`**
- The **`mod`** keyword **defines** a new module as child of the current module
- The **`use`** keyword **imports** a module or item from another module
- Path in module tree defines the file path, with three options:
  - Inline in file of parent module
  - `<module_name>.rs`
  - `<module_name>/mod.rs`

```
// import HashMap from the standard library
use std::collections::hash_map::HashMap;

// define a new module, implementation
// lives in child_module.rs
mod child_module;

// No auto-import! We need to import everything
// we want to use from child_module
use child_module::{MyStruct, my_function};
// equivalent (if this is main.rs)
use crate::child_module::{MyStruct, my_function};
```



# Macros

- Tool for code generation  
⇒ same basic idea as in C/C++

## Not a Footgun

- Unlike C, Rust macros are "hygienic": can not interfere with syntax or local variables
- Operate on the Abstract Syntax Tree  
⇒ only specific syntax is accepted

## Function-like Macros

- Look like functions except for the !
  - `println!` – printing with compile-time checked formatting
  - `dbg!` – quick debug output
  - `panic!`
  - `todo!`
  - `assert!`
  - ...

## Attribute-like Macros

- Usable to annotate any item
- Syntax: `#[attribute_name]`
- Example: auto-implementation of printing, comparisons, etc. for structs via `#[derive(Debug, Eq, ...)]`

# More about the Rust Toolchain

- cargo integrates everything related to building, dependencies and much more
  - Building: `cargo check/build/run [--release]`
  - Dependency management: `cargo add/remove/update <dep>`
  - Tests and Benchmarks: `cargo test/bench`
  - Formatting: `cargo fmt`
  - Linting: `cargo clippy`  
`[- - -W clippy::pedantic]`
- Third-party functionality can be added via `cargo install`
  - Recommendation: `cargo-asm`

## Configuration

- Central configuration file: `Cargo.toml`
  - *`[profile.dev]` and `[profile.release]` for compilation options*
- Integrated option for LTO (link time optimization): `lto = true/"thin"`
  - *Often important, since proper inter-crate inlining requires explicit `#[inline]` annotation*
- No good support for architecture-specific compilation yet (currently: environment variable `RUSTFLAGS="-C target-cpu=native"`)

# Summary

- Rust achieves safety **and** efficiency via the safe/unsafe separation
- No inheritance or metadata: structs are just plain data
- **Pattern matching!**
  - ⇒ Also useful for nullability and error handling
  - ⇒ Rust requires a different approach than typical OO
- Read the **compiler errors**. Completely. Yes, I'm serious!

## Next Lecture

- Ownership and Borrowing
- Generics and Traits
- Is Rust object oriented or functional?