


Rust - Ownership and Borrowing

Colin Bretl, Nikolai Maas, Peter Maucher | 24.10.2024



Recap

1. What is your first impression of Rust?
2. What do you think about pattern matching?

Table of Contents

1. Ownership and Borrowing
2. Basic Ownership
3. References and The Hierarchy of Power
4. Borrow Checker – Motivation
5. Lifetimes
6. Ownership and Borrowing – Intermediate Conclusion
7. Complex Ownership
8. Generics and Traits
9. More on Generics
10. Trait Objects
11. Rust is Object Oriented
12. Rust is Functional
13. So What?

Ownership and Borrowing

developer.okta.com/blog/2022/03/18/programming-security-and-why-rust

- **Topic of the lecture:** understanding how simultaneous safety and efficiency is possible
- Safety = no undefined behavior ↗
- Note: Before Rust, it was not widely accepted that this is **even possible**

Why is Safety Important?

- Computing the **wrong result** fast is not that helpful
- UB is often a security issue
- Avoiding UB is doable, but does not scale well to large teams
⇒ human efficiency!
- Programming without UB is much more enjoyable!*

** author's note: we exclude masochists here for the sake of clarity*

Basic Ownership

Ownership in C++ and Rust

- Without garbage collection, another cleanup mechanism is necessary
- Solution: **RAII** pattern (resource acquisition is initialization)
 - *Single owner for each object: a variable or another object*
 - *Call destructor (in Rust: `drop`) when a variable goes out of scope*

⇒ Typically results in a **tree structure**

Difference: Destructive Moves

- C++: moved-from object still exists in "null" state

```
std::string s1("abc");  
std::print("{} ", s1.size()); // 3  
std::string s2 = std::move(s1);  
std::print("{} ", s1.size()); // 0
```

- Rust: compiler forbids reuse (this also saves a drop call!)

```
let s1 = String::from("abc");  
println!("{}", s1.len()); // 3  
let s2 = s1;  
println!("{}", s1.len());  
// error[E0382]: borrow of moved value: `s1`
```

Basic Ownership

Two Kinds of Objects

- Not all objects hold a resource (e.g., allocation or file descriptor)
- "Plain" structs are marked via the **Copy**-trait
⇒ copy-semantics

COPY

```
let s1: i32 = 42;  
let s2 = s1;  
println!("{}", s1);  
// 42
```

NON-COPY

```
let s1: String = String::from("42");  
let s2 = s1;  
println!("{}", s1);  
// error[E0382]: borrow of moved value: `s1`
```

Aside: Traits

- For now, think **trait** = **interface** (*details later*)
- **Marker traits** without methods mark certain properties of objects

Efficiency

- Moves **and** copies are a bitwise copy of the struct members
 - No unexpected allocations like in C++
 - As efficient as it gets
 - Problematic for self-references

Basic Ownership

Copy and Clone

- **Copy**: struct supports automatic, cheap copies
 - *Edge case: very large structs, e.g. arrays*
- How to create copies of non-copy structs?
- **Clone**: allows to do non-trivial copies via a call to `.clone()`
 - *Cloning typically creates a new allocation*

⇒ New allocations are clearly marked since explicit calls are needed!

Drop

- Non-copy structs often need cleanup code (e.g. freeing the allocation)
- Implemented via the **Drop**-trait
- Automatically called at the end of the current **lexical scope**
- Early drop via `std::mem::drop`

Implementing `std::mem::drop`

```
pub fn drop<T>(_x: T) {}  
// automatic drop at end of scope!
```

Yes, that's everything!

```
let s = String::from("42");  
std::mem::drop(s);  
//           ^^^^^^^ s is moved into the function  
println!("{}", s);  
// error[E0382]: borrow of moved value: `s`
```

References and The Hierarchy of Power

References in Rust

- `&mut`: exclusive reference
(also called "mutable" reference)
- `&`: shared reference
(also called "immutable" reference)
- Unlike C++, references are normal values (e.g. reassignable)
- References have an associated lifetime ↗

Usage of References

- Explicit syntax similar to C pointers
 - `&/&mut` to create a reference (borrowing)
 - `*` for dereferencing

```
fn longest_string(input: &[String]) -> &String {  
    let mut longest: &String = &input[0];  
    for s in &input[1..] { // skip first element  
        if s.len() > longest.len() {  
            longest = s; // reassign reference  
        }  
    }  
    longest  
}
```

(actually, `&str` would make more sense here)

References and The Hierarchy of Power

References and Typing

- References are full-blown types *(with a lifetime)*
- `mut` annotations for variables are **not** part of the type

- `fn foo(mut val: &i32)`

is exactly the same as

```
fn foo(val: &i32)
```

- `fn foo(val: &mut i32)`

is very different to

```
fn foo(val: &i32)
```

Three Modes of Data

- **Owned** (T): arbitrary mutation and destruction
- **Exclusive access** (&mut T): arbitrary mutation, but value must remain valid
- **Shared access** (&T): read-only *(except in case of interior mutability ↗)*

⇒ Each level has strictly more restrictions than previous one

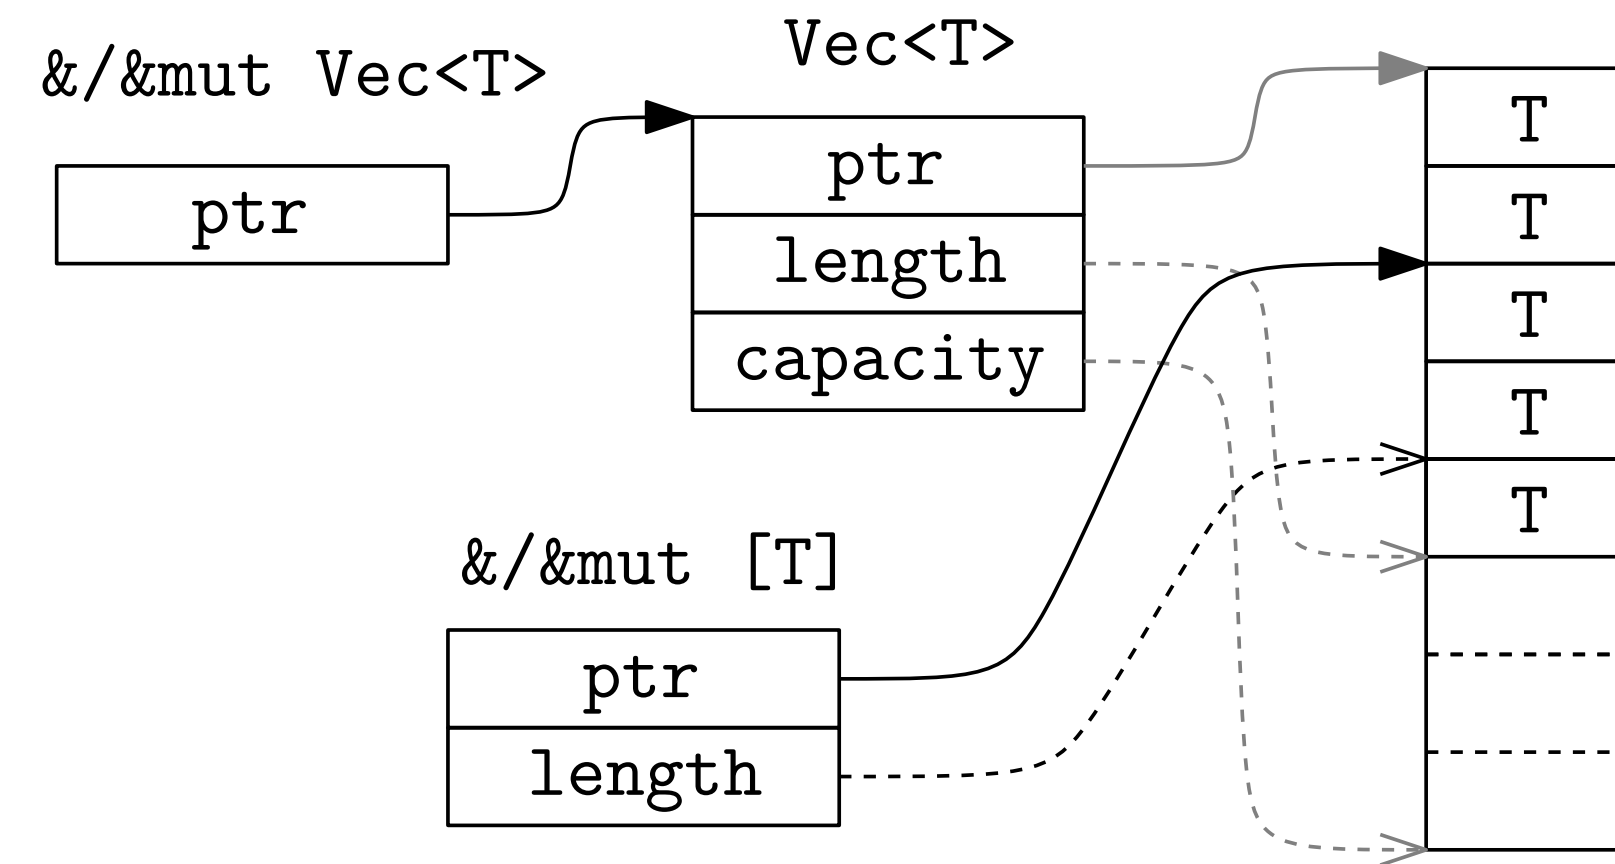
- **Note:** this mostly applies to C++,* too, but Rust makes it more explicit

* let's ignore `const_cast` here

References and The Hierarchy of Power

Owned and Non-Owned Types

- `Vec<T>`: container with owned data
- `&/&mut Vec<T>`: reference to a `Vec<T>`
- `&/&mut [T]`: slice, i.e., reference to a contiguous part of the data
- `[T]`: also a type on its own – data of unknown length (dynamically sized)



Even More Types

- `String` and `&str` work analogous to `Vec<T>` and `&[T]`
- `Box<[T]>`, `Rc<[T]>` and `*const [T]` (smart/raw pointer to `[T]`)
- The `Fn`, `FnMut`, `FnOnce` hierarchy (traits for functions and lambdas)

⇒ Using `&Vec<T>` doesn't make sense,
`&[T]` is superior (compare `&i32/i32`)

- *Though it might appear in generic code*

Borrow Checker – Motivation

The in the Room

- References have a **lifetime** and are subject to **borrow checking**
- $\&T$ is actually $\&'a\ T$, where $'a$ is a (usually anonymous) lifetime
- Borrow checking (severely) restricts the valid usage of references

Borrow checking ensures that ...

- ... no reference outlives its data
- ... exclusive references are actually exclusive

THE GOLDEN RULE OF BORROW CHECKING

There may be **either** one exclusive reference **or** multiple shared references to the same data at any given time.

- Why does this make sense?

Borrow Checker – Motivation

The Ingredients of Safety

- No nullability
- Bounds checking for array accesses, well defined overflow
- **Separating safe and unsafe code** via safe wrapper APIs
 - Safe API = no possible usage can result in undefined behavior (UB)
- The **borrow checker** (?)
- Actually, why do we even need the borrow checker?

"Safe C++"

- **Thought experiment:** apply first three points to C++
 - *Add nullability checks and bounds checks*
 - *Make overflow (and similar "unnecessary" undefined behavior) well defined*
 - *Add `unsafe` keyword to mark remaining UB*
 - Require safe wrappers around any unsafe code
- ⇒ We just made C++ safe*
(or did we?)

* *and we are not the first to try*

Borrow Checker – Motivation

"Safe C++" – the Unsafe Parts

- Main question: What needs to be unsafe?
⇒ Let's look at an example!

```
int push_and_compute(std::vector<int>& input,  
                    const int& val) {  
    int new_element = compute_from(val);  
    input.push_back(new_element);  
    return compute_result(val);  
}
```

```
int main() {  
    std::vector<int> input = {1};  
    return push_and_compute(input, input[0]);  
}
```

- This code is completely fine
- ...
- right?



AddressSanitizer: heap-use-after-free
on address 0x602000000010

Borrow Checker – Motivation

"Safe C++" – the Problems

```
std::vector<int> input = {1};  
return push_and_compute(input, input[0]);
```

- What happened?
 - The vector reallocated its content
 - The reference got **invalidated**
- Source of the problem: the referenced data changed (aka **aliasing**)*

⇒ Most uses of references, pointers
or

iterators must be unsafe

... but this is the majority of C++ code?!

AddressSanitizer: heap-use-after-free
on address 0x602000000010

A Possible Solution

- Let the compiler track the **origin of references** to prevent invalid aliasing
- **Congratulations!**
We just rediscovered the basic idea of the borrow checker

** in a multi-threaded setting, this is problematic even without data
changing its location (data races)*

Borrow Checker – Motivation

- For most data types, mutation requires an exclusive reference
⇒ Prevents referenced data from changing unexpectedly
- References are invalidated if an object is moved or dropped

```
let mut input = vec![1];  
push_and_compute(&mut input, &input[0]);
```

```
error[E0502]: cannot borrow `input` as immutable because it is also borrowed as mutable  
--> src/main.rs:11:35  
11 |   push_and_compute(&mut input, &input[0]);  
   |   ----- ^^^^^ immutable borrow occurs here  
   |   |         |  
   |   |         mutable borrow occurs here  
   |   mutable borrow later used by call
```

There may be **either** one exclusive reference **or** multiple shared references to the same data at any given time.

Lifetimes

- References might be function inputs or outputs
- How does the borrow checker know the origin of returned references?

⇒ It does not, we need to specify!

LOCALITY RULE:

All information relevant for type checking or borrow checking must be contained in the **function signature**.

```
fn which_lifetime(  
    a: &i32, b: &i32, c: &i32  
) -> &i32 {  
    if cond1 { a }  
    else if cond2 { b }  
    else { c }  
}
```

error[E0106]
: missing lifetime specifier

Lifetimes

- Lifetimes describe the **possible origins** of a reference...
- ...or, equivalently, how long it is valid

```
fn which_lifetime<'lt>(  
    a: &'lt i32, b: &'lt i32, c: &'lt i32  
) -> &'lt i32 {  
    if cond1 { a } else if cond2 { b } else { c }  
}
```

"The returned reference might originate from a, b or c."

"The returned reference must not live longer than (any of) a, b or c."

```
fn which_lifetime<'lt>(  
    a: &'lt i32, b: &i32, c: &i32  
) -> Option<&'lt i32> {  
    if cond(b, c) { Some(a) } else { None }  
}
```

"The returned reference always originates from a."

⇒ Lifetimes allow to track reference origins using only function signatures*

** yes, the syntax is a bit ... special*

Lifetimes

Lifetime Elision

- Only lifetimes relevant for the output need to be specified
- Often, lifetimes can be fully elided (*i.e. inferred by the compiler using fixed rules*)
 - Only one input lifetime
 - Methods: lifetime originates from **&self**

```
fn which_lifetime<'lt>(  
    a: &'lt i32, b: &i32, c: &i32  
) -> Option<&'lt i32> {  
    if cond(b, c) { Some(a) } else { None }  
}
```

implicitly expands to

```
fn which_lifetime<'lt, 'b, 'c>(  
    a: &'lt i32, b: &'b i32, c: &'c i32  
) -> Option<&'lt i32> {  
    if cond(b, c) { Some(a) } else { None }  
}
```

Rule of thumb: one named lifetime suffices in 95% of cases

Lifetimes

Lifetimes in Structs

```
struct ContainsReference<'ref> {  
    ref_data: &'ref i32,  
}
```

- This compiles...
- ...but is a bad implementation!

```
fn transform(val: &i32) -> &i32 {  
    ContainsReference {  
        ref_data: val  
    }.get_data()  
} // error: cannot return value  
  // referencing temporary value
```

```
impl<'ref> ContainsReference<'ref> {  
    fn get_data(&self) -> &i32 {  
        self.ref_data  
    }  
}
```

expands to

```
fn get_data<'self>(&'self self) -> &'self i32 {  
    self.ref_data  
}
```

- But `&self` is not the actual origin!

Fixed version:

```
impl<'ref> ContainsReference<'ref> {  
    fn get_data(&self) -> &'ref i32 {  
        self.ref_data  
    }  
}
```

Ownership and Borrowing – Intermediate Conclusion





- Ensure memory safety by preventing aliasing
- The problem: combining sharing and mutation

More than Memory Safety

```
fn some_business_logic(data: &mut DataHolder,
                       value: &Value) {
    let old_val = *value;
    some_subroutine(data);
    assert_eq!(old_val, *value);
    // Rust guarantees this always holds
}
```

- Typical OO programs have no such guarantee (*data might reference value*)
- Rust makes mutation **locally** explicit
- Aliasing can also cause logic bugs!

Locality of Reasoning

	No mutation	Mutation
No sharing		
Sharing		

- Common adage in Rust: "if it compiles, it works"*

* except, of course, for logic bugs not related to aliasing, as well as memory leaks, deadlocks, ...

Ownership and Borrowing – Intermediate Conclusion

- Single owner principle leads to tree-shaped ownership
- Referenced data may not be mutated through other paths

⇒ Storing references permanently is usually a **bad idea** (*Generally, references should be kept "stack-downwards" of the origin*)

Benefits

- Safety! Correctness!
 - Lifetimes allow complex computations with references
 - We can use references "fearlessly"
- ⇒ Avoids defensive copies

Drawbacks

- Enforces rather restrictive data layout
- Manually specified lifetimes are tricky (lifetime errors...)

What if the data is inherently not tree-shaped?

Complex Ownership

Strategy 1: Smart Pointers and Interior Mutability

- Smart pointers support different ownership strategies

C++

`std::unique_ptr<T>`

`std::shared_ptr<T>`

Rust

`Box<T>`

`Rc<T>` *(not thread-safe)*

`Arc<T>` *(thread-safe)*

Heap-allocated owned data

Reference counted, shared ownership


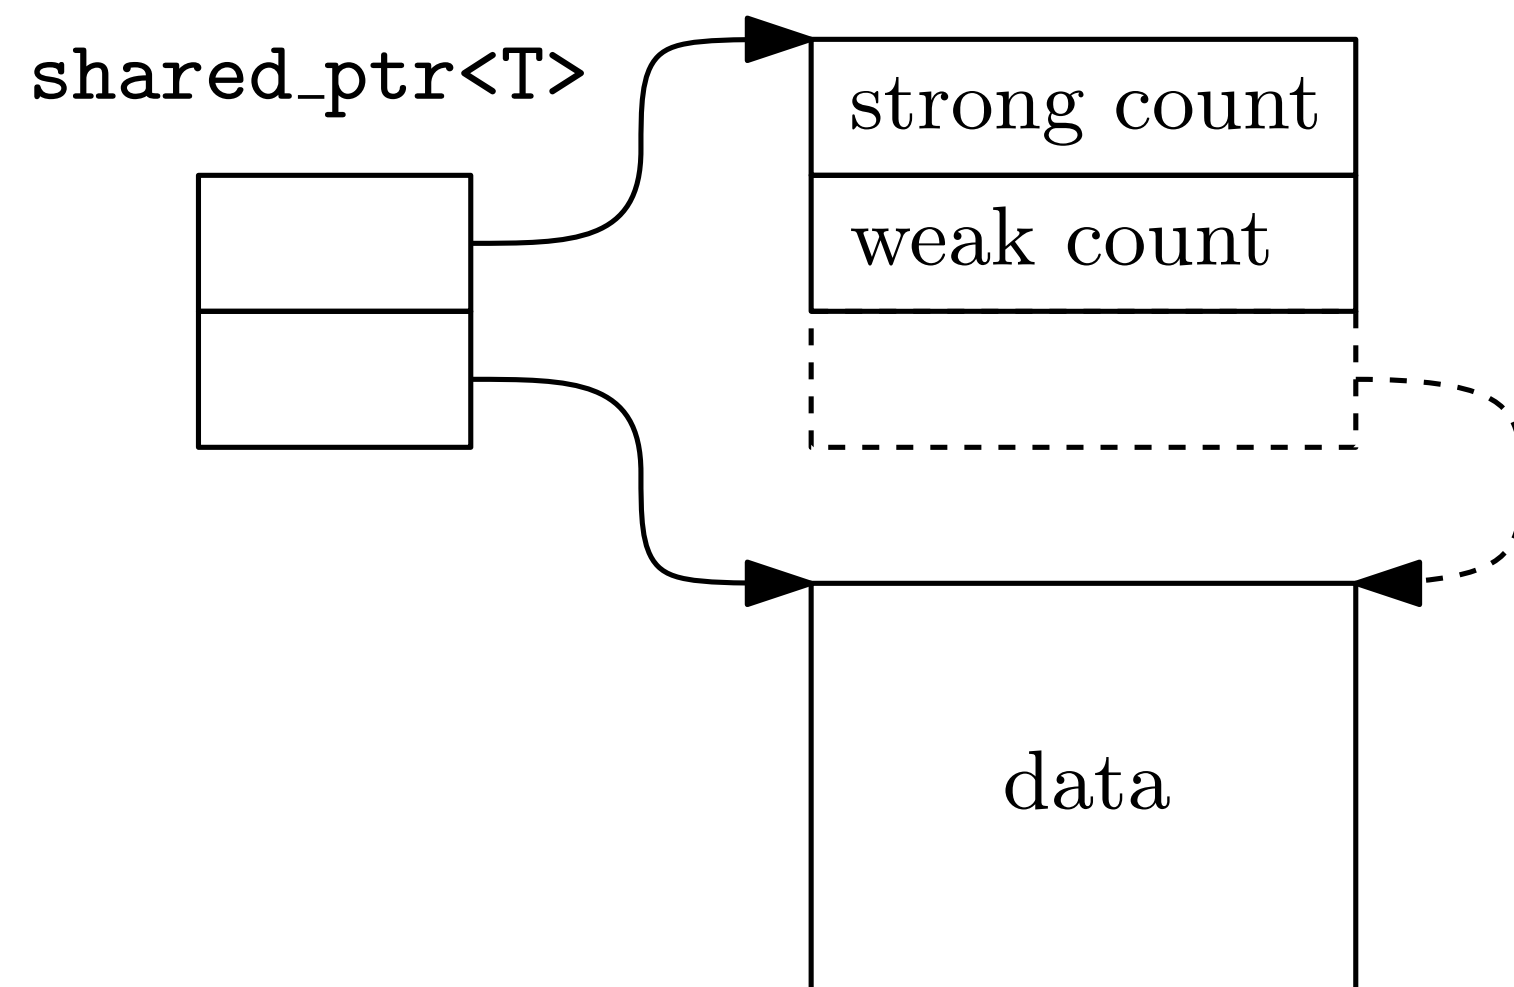
Atomic reference counted, shared ownership

- Careful: reference counted **cycles** can cause memory leak!

Smart Pointers and Interior Mutability

- Rc allows single-threaded reference counting without synchronization overhead
 - *Not available in C++!*

- Rust always uses the same allocation for counter and data
 - 👍 Better memory layout
 - 👎 Creation requires copying the data



Smart Pointers and Interior Mutability

- **Restriction:** only allows shared (= immutable?) references to the data
- Solution: Interior Mutability!

Interior Mutability: pattern that allows mutation through shared references by using specialized container types

TYPES WITH INTERIOR MUTABILITY

<code>RefCell<T></code>	checks at runtime for exclusiveness, returning an exclusive reference if possible
<code>Cell<T></code>	allows mutation only via copies (no reference involved)
<code>Mutex<T></code>	ensures thread-safe exclusive access via locking
<code>UnsafeCell<T></code>	unsafe core primitive for interior mutability, basis for all the other types

Smart Pointers and Interior Mutability

Example: Updating Data Inside a Mutex

```
let data: Arc<Mutex<MyData>> = initial_arc.clone();  
// ...  
let mut guard = data.lock().unwrap(); // lock the mutex  
    // ^-- the guard ensures exclusive access  
*guard = computeUpdate(guard.borrow());  
mem::drop(guard); // unlock the mutex  
// ...
```

Smart Pointers and Interior Mutability

- Reference counting
 - Efficient access
 - Memory overhead: two counters, separate allocation
- Interior Mutability
 - Efficient for small values via `Cell`, but copying inefficient for large data
 - `RefCell` needs runtime-check

Benefits

- Allows safe shared ownership
- Reference counting for large immutable data chunks is efficient
- Enables some typical OO patterns (objects referencing each other)

Drawbacks

- Significant memory overhead if small data chunks are shared
- Mutating large data via `RefCell` involves runtime-check
- OO patterns are often a bad idea (*in Rust*)

⇒ Design your data layout to avoid interior mutability,
but use it where necessary

Complex Ownership

Strategy 2: Arena Allocation

- Pattern for handling many identical objects
- Idea: don't use pointers, instead use a central data structure (= arena) and IDs
- Simple variant: `Vec<T>` with indices as IDs

Avoiding the borrow checker

- No borrow checking on IDs \Rightarrow no issues with sharing
- Borrow checking now happens on the arena reference
- If elements can be deleted: **ID invalidation**
(*"safe" version of use-after-free*)
 - Can be mitigated with runtime checks \Rightarrow overhead!
(*see: generational indices, the slotmap library*)

Example: A Static Graph Data Structure

```
// macro to derive basic properties for IDs
#[derive(Debug, Clone, Copy, ...)]
pub struct NodeID(pub u32);
#[derive(Debug, Clone, Copy, ...)]
pub struct EdgeID(pub u32);

struct NodeData{ first_edge: EdgeID, weight: i32 }
struct EdgeData{ target_node: NodeID, weight: i32 }

pub struct Graph { nodes: Vec<NodeData>, edges: Vec<EdgeData> }

impl Graph {
    pub fn node_weight(&self, node: NodeID) -> i32 {
        self.nodes[node.0 as usize].weight
    }

    pub fn nodes(&self) -> impl Iterator<Item=NodeID> {
        // Iterator over nodes. Remove last entry (which is a sentinel)
        (0..(self.nodes.len() - 1) as u32).map(NodeID)
    }
}
```


Arena Allocation

- Safety: access via the ID needs a bounds check
⇒ Provide an unsafe getter for hot code

Benefits

- Elegantly avoids borrow checker issues
- Very efficient memory layout and allocation of elements
- For smaller data sets, we can use 32-bit IDs

Drawbacks

- Avoiding the borrow checker means less safety: ID invalidation
- Always requires a reference to the arena
- Only with bounds checking safe
- Compared to pointers, access requires an additional offset calculation

⇒ Always consider using arena allocation
if there are many similar objects!

Complex Ownership

Strategy 3: **unsafe** and Raw Pointers

- Sometimes we can't accept any overhead
⇒ This is the raison d'être for **unsafe**
- Rust supports raw pointers that are mostly equivalent to C
(*more details in the next lecture ↗*)

Benefits

- No overhead
- Can be encapsulated in a safe data structure
- Sometimes there is **no other possibility**

Drawbacks

- Opens the door to undefined behavior
- Some rules are even trickier than C/C++

Complex Ownership

Strategy 1: Smart Pointers and Interior Mutability

- Enables shared ownership and mutation at once
- Has **inherent overhead** and is not very ergonomic to use

Strategy 2: Arena Allocation

- A generally useful and quite efficient strategy
- Not always applicable (needs many similar objects)

Strategy 3: **unsafe** and Raw Pointers

- Useful if no overhead is acceptable
- Allows efficient custom data structures

Note: Arena allocation is often **more efficient** 🚴 than raw pointers!

Generics and Traits

- Generic programming is necessary for strong abstraction support
- "Safer" variant of templates: generic parameters must be annotated with **trait bounds** (*similar to enforced concepts*)
- Implemented via monomorphization (separate compilation of each variant)

```
struct GenericData<T> {  
    data: Vec<T>  
}  
  
impl<T> GenericData<T> {  
    // trait bound: S must be convertible to T  
    fn add<S: Into<T>>(&mut self, value: S) {  
        self.data.push(value.into());  
    }  
}
```

Traits

- Traits define shared functionality of different types
- Includes constants, types, static functions and methods

```
trait AbstractVector {  
    // associated constant  
    const DIMENSION: usize;  
    // associated type  
    type Scalar;  
    // associated function  
    fn new_zeroed() -> Self;  
    // associated method  
    fn get(&self, index: usize)  
        -> Option<Self::Scalar>;  
}
```

Generics and Traits

Using Traits

- Generic code must specify all required properties via trait bounds
- `where` allows to conveniently add more complex bounds

```
fn compute_sum<V: AbstractVector>(value: &V)
    -> V::Scalar
    // additional bounds
    where V::Scalar: AddAssign + Default
{
    // Default trait for initial zero
    let mut result = Default::default();
    for index in 0..V::DIMENSION {
        // AddAssign trait for addition
        result += value.get(index).unwrap();
    }
    result
}
```

```
trait AbstractVector {
    const DIMENSION: usize;
    type Scalar;
    fn new_zeroed() -> Self;
    fn get(&self, index: usize)
        -> Option<Self::Scalar>;
}
```

Generics and Traits

Implementing Traits

- Traits are implemented via a specific `impl` block
- Traits can be implemented for **already existing** types! (*aka extension methods*)

```
impl AbstractVector for f64 {  
    const DIMENSION: usize = 1;  
    type Scalar = f64;  
  
    fn new_zeroed() -> Self { 0.0 }  
  
    fn get(&self, index: usize) -> Option<f64> {  
        (index == 0).then_some(*self)  
    }  
}
```

```
trait AbstractVector {  
    const DIMENSION: usize;  
    type Scalar;  
    fn new_zeroed() -> Self;  
    fn get(&self, index: usize)  
        -> Option<Self::Scalar>;  
}
```

Orphan rule (simplified):

A trait implementation is valid if either the trait or the implementing type is local.

More on Generics

- **Three kinds** of generic parameters:
 - Lifetimes*
 - Type parameters
 - Const generics

```
struct RefToArray<'a, Content,  
                const N: usize> {  
    data: &'a [Content; N],  
}
```

Similarities to C++

- Type parameters + traits are **turing complete**
- Compilation of generics via **monomorphization**
⇒ Zero overhead! 📌
- **const** is like **constexpr**, but more restricted (for now)

Differences to C++

- Type checking happens **before** monomorphization
 - Enabled by the explicit bounds
⇒ Errors are detected earlier (at declaration instead of usage)
⇒ Better error messages
- Explicit bounds can be restrictive

** lifetimes are a bit special: they are erased instead of monomorphized and allow subtyping through variance*

More on Generics

```
fn try_create() {  
    let mut vec = Vec::<Uncloneable>::new();  
    vec.resize(10, Uncloneable);  
}
```

```
void try_create() {  
    std::vector<Uncopyable> vec;  
    vec.resize(10, Uncopyable{});  
}
```

```
error[E0599]: the method `resize` exists for struct `Vec<Uncloneable>`,  
but its trait bounds were not satisfied  
--> <source>:7:9  
|  
2 | struct Uncloneable;  
| ----- doesn't satisfy `Uncloneable: Clone`  
...  
7 |     vec.resize(10, Uncloneable);  
|     ^^^^^^^  
|  
= note: the following trait bounds were not satisfied:  
    `Uncloneable: Clone`  
help: consider annotating `Uncloneable` with `#[derive(Clone)]`  
|  
2 + #[derive(Clone)]  
3 | struct Uncloneable;  
|  
error: aborting due to 1 previous error
```

```
In file included from <source>:2:  
In file included from /opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/vector:63:  
In file included from /opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/allocator.h:46:  
In file included from /opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/x86_64-linux-gnu/bits/c++allocator.h:33:  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/new_allocator.h:187:23: error: call to deleted  
constructor of 'Uncopyable'  
187 |     { ::new((void *)__p) _Up(std::forward<Args>(__args)...); }  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/alloc_traits.h:537:8: note: in instantiation  
of function template specialization 'std::_new_allocator<Uncopyable>::construct<Uncopyable, const Uncopyable &>' requested here  
537 |     __a.construct(__p, std::forward<Args>(__args)...);  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_vector.h:1832:21: note: in instantiation  
of function template specialization 'std::allocator_traits<std::allocator<Uncopyable>>::construct<Uncopyable, const Uncopyable &>' requested here  
1832 |     _Alloc_traits::construct(_M_this->_M_impl, _M_ptr(),  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/vector.tcc:540:25: note: in instantiation of  
function template specialization 'std::vector<Uncopyable>::_Temporary_value::_Temporary_value<const Uncopyable &>' requested here  
540 |     _Temporary_value __tmp(this, __x);  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_vector.h:1034:4: note: in instantiation of  
member function 'std::vector<Uncopyable>::_M_fill_insert' requested here  
1034 |     _M_fill_insert(end(), __new_size - size(), __x);  
|     ^  
<source>:15:9: note: in instantiation of member function 'std::vector<Uncopyable>::resize' requested here  
15 |     vec.resize(10, Uncopyable{});  
|     ^  
<source>:10:5: note: 'Uncopyable' has been explicitly marked deleted here  
10 |     Uncopyable(const Uncopyable& other) = delete;  
|     ^  
In file included from <source>:2:  
In file included from /opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/vector:65:  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_uninitialized.h:90:21: error: static  
assertion failed due to requirement 'is_constructible<Uncopyable, Uncopyable &&::value>': result type must be constructible from input type  
90 |     static_assert(is_constructible<_ValueType, _Tp>::value,  
|                   ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_uninitialized.h:182:4: note: in  
instantiation of function template specialization 'std::__check_constructible<Uncopyable, Uncopyable &&::value>' requested here  
182 |     = _GLIBCXX_USE_ASSIGN_FOR_INIT(_ValueType2, _From);  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_uninitialized.h:101:13: note: expanded  
from macro '_GLIBCXX_USE_ASSIGN_FOR_INIT'  
101 |     && std::__check_constructible<T, U>()  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_uninitialized.h:373:19: note: in  
instantiation of function template specialization 'std::uninitialized_copy<std::move_iterator<Uncopyable >, Uncopyable >' requested here  
373 |     return std::uninitialized_copy(__first, __last, __result);  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_uninitialized.h:384:19: note: in  
instantiation of function template specialization 'std::__uninitialized_copy_a<std::move_iterator<Uncopyable >, Uncopyable *, Uncopyable>' requested here  
384 |     return std::__uninitialized_copy_a(_GLIBCXX_MAKE_MOVE_ITERATOR(__first),  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/vector.tcc:548:10: note: in instantiation of  
function template specialization 'std::__uninitialized_move_a<Uncopyable *, std::allocator<Uncopyable>>' requested here  
548 |     std::__uninitialized_move_a(__old_finish - __n,  
|     ^  
/opt/compiler-explorer/gcc-13.2.0/lib/gcc/x86_64-linux-gnu/13.2.0/../../../../include/c++/13.2.0/bits/stl_vector.h:1034:4: note: in instantiation of  
member function 'std::vector<Uncopyable>::_M_fill_insert' requested here  
1034 |     _M_fill_insert(end(), __new_size - size(), __x);  
|     ^
```

** this is the clang error which is IMO notably better than the gcc error: at least the actual problem is stated at the beginning*

Trait Objects

- Sometimes, we really need **dynamic dispatch**
 - Avoid code bloat
 - Handling different types uniformly at runtime (type erasure)
 - Example: Library that supports user-defined types
- Dynamic dispatch works via traits, too
 - *Note: C++ virtual methods are a completely separate system from templates*
- **Trait objects** are marked with **dyn** and must be behind some kind of pointer


```
trait AbstractVectorObj {  
    type Scalar;  
    fn get(&self, index: usize)  
        -> Option<Self::Scalar>;  
}
```

```
let mut list = Vec::new();  
let first: Box<dyn AbstractVectorObj<  
    Scalar = f64>> = Box::new(1.0);  
list.push(first);  
let second: Box<dyn AbstractVectorObj<  
    Scalar = f64>> = Box::new([24.0, 42.0]);  
list.push(second);  
let scalar: f64 = list[0].get(0);
```

Trait Objects

- Based on vtables as in C++
- vtable-pointer **not** part of the struct (*structs are plain data!*)
- **Fat pointers:** pointer to the trait object also contains the vtable-pointer

⇒ Restricted to functionality that can be implemented with a vtable!



Object Safety

- No associated constants
- All functions must be methods*
- No generic functions (except for lifetimes)
- Associated types must be specified*

* *functions and associated types can be "excluded" from trait objects via **where Self: Sized***

Trait Objects

- Different memory layout than C++
- vtable-pointer is created "on the fly"

Benefits

- Unified system for static and dynamic polymorphism
- Zero overhead when not used

Drawbacks

- Object safety can be awkward
- Many fat pointers to same object create memory overhead

Rust is Object Oriented

Support for many OO features

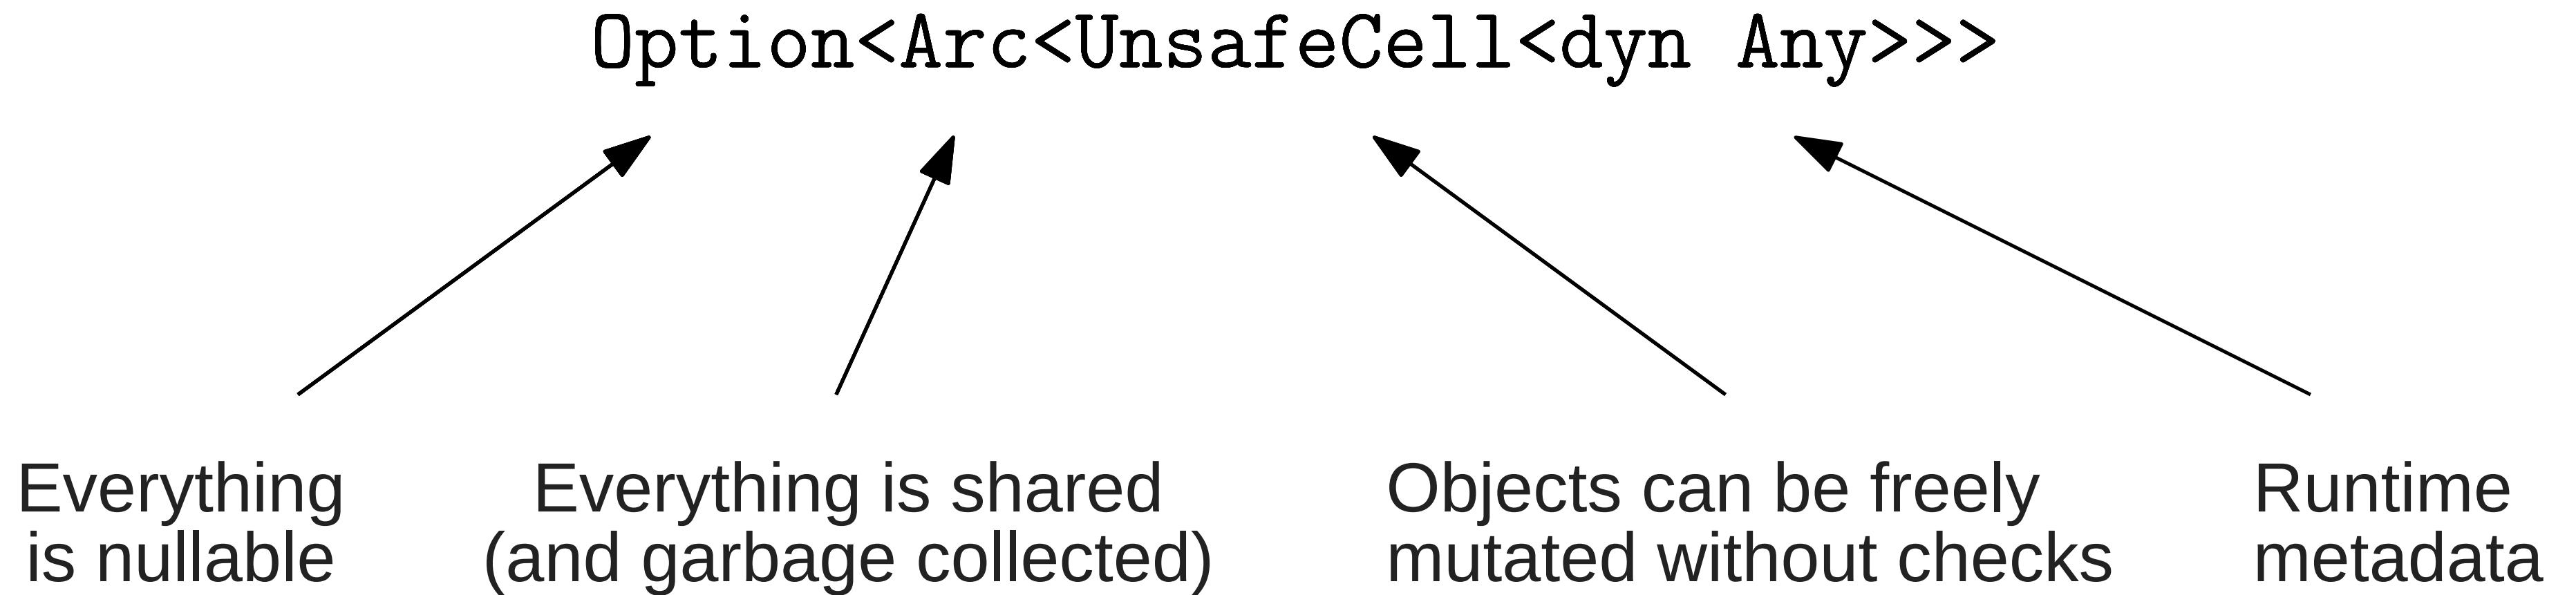
- Methods allow to use structs similar to objects
- Structs can contain data and behavior
- Dynamic dispatch via trait objects
- Traits can "inherit" from other traits

Encapsulation

- Visibility management via `pub`
- Hiding of implementation details and preserving invariants
- Traits for shared behavior

Rust is Not Object Oriented

The equivalent of a normal Java object in Rust:



- Rust has much more restrictions on data layout
 - Shared ownership is heavily discouraged
- ⇒ Typical OO "object soup" architecture is very unidiomatic (and unergonomic)

Rust is Functional

Support for many functional features

- Immutability by default
- Pattern matching
- Iterators with `filter`, `map`, `fold`, etc.
- Monadic types
(`Option`, `Result`, `Iterator`, ...)

Mutability and Type Safety

- Similar to the functional paradigm, Rust restricts mutability and avoids it where unnecessary
- A strong type system for catching as many errors at compile time as possible

Functional Programming:

A programming paradigm that focuses on avoiding state and mutability and uses pure functions as fundamental abstraction.

Rust is Not Functional

- Functional languages tend to avoid mutability completely
- Requires according data structures and optimizations to be approximately efficient
⇒ Makes many common algorithms impossible
- Rust restricts but still embraces mutability
- Rust has (intentionally) no immutable data structures in the standard library

"[...] pure functional programming is an ingenious trick to show you can code without mutation, but Rust is an even cleverer trick to show you can just have mutation."*

* without.boats/blog/notes-on-a-smaller-rust/

So What?

Conclusion

- Rust combines many of the best features of OO and functional...
- ...and comes with new, previously unknown, drawbacks (**lifetimes!**)
- While Rust has **features** from multiple paradigms, both OO and functional **program structure** are a bad fit

Structure of Rust Programs

- Rust programs tend to use a simple procedural structure
 - Algorithms are implemented within (generic) free functions
 - Explicit flat data layout using arenas
 - Types encode invariants of the data

⇒ A flavor of procedural? A new paradigm?*

- *How about ownership oriented programming (OOP)? ... oh wait*

* *at least, Rust has inspired some others*